

Neural-Symbolic VideoQA: Learning Compositional Spatio-Temporal Reasoning for Real-world Video Question Answering

Lili Liang, Guanglu Sun, Jin Qiu, Lizhong Zhang

Harbin University of Science and Technology, School of Computer Science and Technology
Harbin, 157000

Abstract

Compositional spatio-temporal reasoning poses a significant challenge in the field of video question answering (VideoQA). Existing approaches struggle to establish effective symbolic reasoning structures, which are crucial for answering compositional spatio-temporal questions. To address this challenge, we propose a neural-symbolic framework called Neural-Symbolic VideoQA (NS-VideoQA), specifically designed for real-world VideoQA tasks. The uniqueness and superiority of NS-VideoQA are two-fold: 1) It proposes a Scene Parser Network (SPN) to transform static-dynamic video scenes into Symbolic Representation (SR), structuralizing persons, objects, relations, and action chronologies. 2) A Symbolic Reasoning Machine (SRM) is designed for top-down question decompositions and bottom-up compositional reasonings. Specifically, a polymorphic program executor is constructed for internally consistent reasoning from SR to the final answer. As a result, Our NS-VideoQA not only improves the compositional spatio-temporal reasoning in real-world VideoQA task, but also enables step-by-step error analysis by tracing the intermediate results. Experimental evaluations on the AGQA Decomp benchmark demonstrate the effectiveness of the proposed NS-VideoQA framework. Empirical studies further confirm that NS-VideoQA exhibits internal consistency in answering compositional questions and significantly improves the capability of spatio-temporal and logical inference for VideoQA tasks.

Keywords— Video Question Answering; Neuro-Symbolic Learning; Compositional Spatio-Temporal Reasoning; Disentangled Representation; Scene Analysis and Understanding

1. Introduction

VideoQA in real-world scenarios is an attractive yet challenging task, as it requires multi-modal understanding and reasoning [1, 2]. A variety of models have achieved

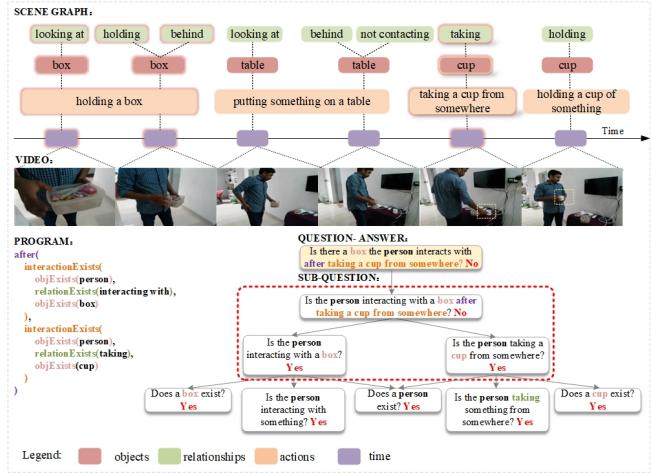


Figure 1. An example of compositional spatio-temporal VideoQA. We mark the objects, relations, actions and time in red, green, orange, and purple respectively. The red dashed box shows a decomposing step from a compositional question to its sub-questions.

state-of-the-art performance by cross-modal alignment [3–6] and transformer-based multi-step reasoning [7–9]. However, these models face challenges when the questions involve complex compositional structures or require understanding spatio-temporal information underlying in videos.

We argue that addressing compositional spatio-temporal questions hinges on two key factors: 1) the cognition of concepts encompassing persons, objects, relations, action chronologies, which contributes to spatio-temporal reasoning; and 2) the compositional reasoning with the analysis of question structures, which effectively mitigates error accumulation during multi-step reasoning.

Using Figure 1 as an example, answering the question “Is there a box the person interacts with after taking a cup from somewhere” requires the ability to recognize the concepts such as the person, objects (box, cup), relations (looking at, taking), and actions (person taking a cup from somewhere) from the given video. Additionally, the question needs to be decomposed into sub-questions, such as “Is the person interacting with a box?” and “Is the person taking a cup from somewhere?”

somewhere?”, which may further require recursive decomposition. Lastly, the ability to reason about the temporal relationship indicated by “after” using appropriate reasoning rules is crucial. For instance, both sub-questions in Figure 1 receive positive answers, but none of the actions after “taking a cup from somewhere” involve interacting with a box, the final answer would be “No”.

Several approaches have made remarkable progress in understanding and reasoning in visual question answering (VisualQA) [10, 11] and VideoQA by incorporating neural-symbolic methods. That is, to decouple object attributes with a neural model and to deduce the answer through symbolic reasoning [12, 13]. However, these approaches are oriented towards synthesized scenes containing geometric solids merely, lacks suitability for real-world videos.

To address this limitation, we propose Neural-Symbolic VideoQA (NS-VideoQA), a framework for real-world compositional spatio-temporal VideoQA, designed with neural cognition and symbolic reasoning. For neural cognition, we introduce a Scene Parser Network (SPN) for static and dynamic scenes in videos. SPN transforms non-structural video data into Symbolic Representation (SR), structuralizing persons, objects, relations, and action chronologies. For symbolic reasoning, we proposed Symbolic Reasoning Machine (SRM), where the top-down question decompositions and bottom-up compositional reasonings are conducted. Specifically, we observed that the reasoning rules varies conditioning by the categories of sub-questions. Therefore, a polymorphic program executor is constructed to reason iteratively with internal consistency in various situations. Our NS-VideoQA not only improves the rationality of compositional spatio-temporal reasoning in real-world VideoQA, but also enabled step-by-step error analysis by tracing the intermediate results along the entire reasoning process.

We evaluate our NS-VideoQA framework on the AGQA Decomp dataset [14], which is focused on compositional spatio-temporal questions. We find that our framework outperforms existing purely neural VideoQA models. Further empirical analysis measuring Compositional Accuracy (CA), Right for the Wrong Reasons (RWR) and Internal Consistency (IC) shows that the NS-VideoQA provides superior capability in compositional spatio-temporal reasoning.

In summary, our contributions are:

1. NS-VideoQA is proposed, which is a neural-symbolic framework for VideoQA that enables compositional spatio-temporal reasoning in real-world videos. This is achieved by transforming the video into symbolic representation and conducting iterative reasoning.
2. Scene Parser Network, a cognitive model based on Transformer, aims to extract symbolic representations from static-dynamic scenes in videos. This approach converts unstructured videos and questions into struc-

- tured representations, facilitating symbolic reasoning.
- 3. Symbolic Reasoning Machine, an polymorphic reasoning engine for compositional answer deduction with human-readable inference traces.

2. Related Work

2.1. Transformer-based VideoQA

In recent years, predicting answers via visual-textual alignment has been the mainstream methodology towards VideoQA. Typical methods align different modalities via cross-attention [15], and conduct multi-step reasoning using multi-hop attention [16, 17] or stacked self-attention layers [18].

Recently, pre-trained Transformers improved both cognition and reasoning in VideoQA. For better cognition, several methods [19, 20] integrate tokens from different modalities and pass them through self-attention layers to enhance attention-based cross-modal cognition. However, these methods are insensitive to temporal dependencies required by questions, resulting in a lack of temporal reasoning. To enhance temporal modeling, All in One [7], PMT [21], and RTransformer [22] focus on capturing and leveraging temporal information to improve the understanding and reasoning capabilities of the models. For better reasoning, MIST [8] and VGT [9] utilize different strategies, such as multi-step spatio-temporal reasoning and graph-based reasoning, to capture the underlying structure in video. However, these models struggle when the questions require multi-step reasoning, which reveals the limitation of purely neural models.

Unlike the approaches mentioned above, we utilize a Transformer-based model solely for recognizing symbolic representations, all compositional spatio-temporal reasoning is performed by algorithms based on symbolic logic.

2.2. Neuro-symbolic Methods

Neural-symbolic reasoning strives to build a more transparent, precise, and interpretable automated reasoning process by integrating symbolic logic with the powerful perceptual capabilities of neural networks [10, 11]. As a pioneering work, IEP [23] applied neural-symbolic reasoning to VisualQA by generating programs from questions and conducting inference using CNN features. More recently, several work [12, 24, 25] improve the explainability of IEP-based VisualQA.

VideoQA presents a more challenging task that requires a symbolic reasoning system to model the temporal and dynamic information in videos. CLEVRER [13], DCL [26], and COMPHY [27] are all approaches that aim to enhance symbolic reasoning in VideoQA by employing different strategies, such as task decomposition, concept grounding, and understanding of dynamics and compositionality.

However, the models were designed for synthesized videos, not applicable for real-world VideoQA with complex visual scenes.

In this work, we expand the applicability of compositional spatio-temporal VideoQA onto real-world videos.

2.3. Compositionality Benchmarks

Recent VideoQA benchmarks [28, 29] indicate that state-of-the-art models struggle to answer compositional questions due to multiple error sources. AGQA [30] and its balanced version AGQA 2.0 [31] was proposed as benchmarks of compositional spatio-temporal reasons for VideoQA. For clear analysis of mispredicting compositional reasoning types, a new benchmark AGQA Decomp [14] was constructed with question programs more relevant to video, providing an systematical environment for compositional reasoning and error analysis. Thus, we adopt AGQA 2.0 to evaluate our NS-VideoQA framework.

3. NS-VideoQA

To enhance the ability of neural-symbolic models in answering compositional spatio-temporal questions about real-world videos, we propose the **Symbolic Representation (SR) of static-dynamic scene** as a symbolic representation of videos. Based on the idea of neural-symbolic reasoning, the NS-VideoQA framework can be divided into two phases.

In the cognition phase, we propose a **static-dynamic scene parser network (SPN)**. To obtain the SR for a given video, a Static Scene Parser detects persons, objects, and relationships in video frames(Figure 2-I). Concurrently, a Dynamic Scene Parser detects action chronologies in video clips(Figure 2-II).

In the symbolic reasoning phase, we propose a **symbolic reasoning machine (SRM)**. To generate the answer for a given question, a Language Question Parser transforms the question from natural language into a symbolic program (Figure 2-III). Subsequently, a Program Symbolic Executor performs the program based on SR (Figure 2-IV).

The NS-VideoQA framework utilizes SPN to convert an input video into SR. It further utilizes SRM to decompose the compositional question into a program. The program, together with reasoning rules, is iteratively applied to SR, ultimately generating the answer. Further details are described below.

3.1. Symbolic Representation

Symbolic representation of static scene. For a video with T frames $v = \{I_1, I_2, \dots, I_T\}$, the Static Scene Parser φ parses relationships between the subject-object pairs in video $\varphi(v) = \{e_1, e_2, \dots, e_N\}$, where $e_i = (s_{sbj}, s_{rel}, s_{obj})$, $i = 1, \dots, N$, s_{sbj} , s_{rel} and s_{obj} indicate subject, relation and object, respectively. The inference model can effectively filter objects based on the symbolic representation of the relationships in question. The symbolic representation of static scene is defined as $R_{static}(v) = \varphi(v)$.

Symbolic representation of dynamic scene. For a video with T clips $v = \{c_1, c_2, \dots, c_T\}$, the Dynamic Scene Parser ψ generates actions in video $\psi(v) = \{y_1, \dots, y_N\}$, where $y_i = (s_{sbj}, s_{act}, s_{obj}, t_{start}, t_{end})$, $i = 1, \dots, N$, s_{sbj} , s_{act} , s_{obj} , t_{start} and t_{end} indicate subject, action, object, starting and ending times, respectively. The reasoning model can effectively filter action instances based on the symbolic representation of the actions, it also can accurately determine their temporal order based on the starting and ending times of the actions. The symbolic representation of dynamic scene is defined as $R_{dynamic}(v) = \psi(v)$.

Therefore, the SR can be expressed as follows:

$$SR(v) = (R_{static}(v), R_{dynamic}(v)) \quad (1)$$

3.2. Static-dynamic Scene Parser Network

Static Scene Parser φ . Inspired by STTran [32], the t -th frame I_t is fed into Object-Relation Detector [33] to detect objects o_t^1, \dots, o_t^N , where o_t^i indicates the i -th object, v_t^i , c_t^i and r_t^i indicate o_t^i corresponding object, semantic and relation feature, respectively.

Relationship representation. It is generated to capture the relationships of subject-object pairs by leveraging object, semantic, and relation features. We choose an object with higher confidence to be “person”, suppose it to be o_t^1 , its relationship with other objects can be expressed:

$$x_t^i = \left(W_s v_t^1, W_o v_t^i, W_r f_r, c_t^1, c_t^i \right), i = 2, \dots, N \quad (2)$$

Where $f_r = \begin{bmatrix} r_t^1 & \dots & r_t^i \end{bmatrix}$ is the concatenation of relation features of o_t^1 and o_t^i , W_s, W_o, W_r are the linear matrices for dimension compression. The relationships of the t -th frame are represented as $X_t = \{x_t^2, \dots, x_t^N\}$.

Encoder. In order to encode the spatial relationship representation of subject-object pairs, an L -layer Transformer encoder [34] is utilized. The query Q , key K , and value V in the encoder share the same input, and the output of the n -layer encoder is:

$$\begin{aligned} X_t^{(n)} &= \text{Att}_{enc} \left(Q = K = V = X_t^{(n-1)} \right) \\ X_t^{(0)} &= X_t \end{aligned} \quad (3)$$

Decoder. A sliding window of size η runs over the sequence of spatial relationship representations $\{X_1^{(L)}, \dots, X_T^{(L)}\}$, generating $Z_j = \{X_j^{(L)}, \dots, X_{j+\eta-1}^{(L)}\}$, $j = 1, \dots, T - \eta + 1$. The decoder takes Z_j as input, and consists of L stacked identical self-attention layers $\text{Att}_{dec}()$. The decoder can be described as

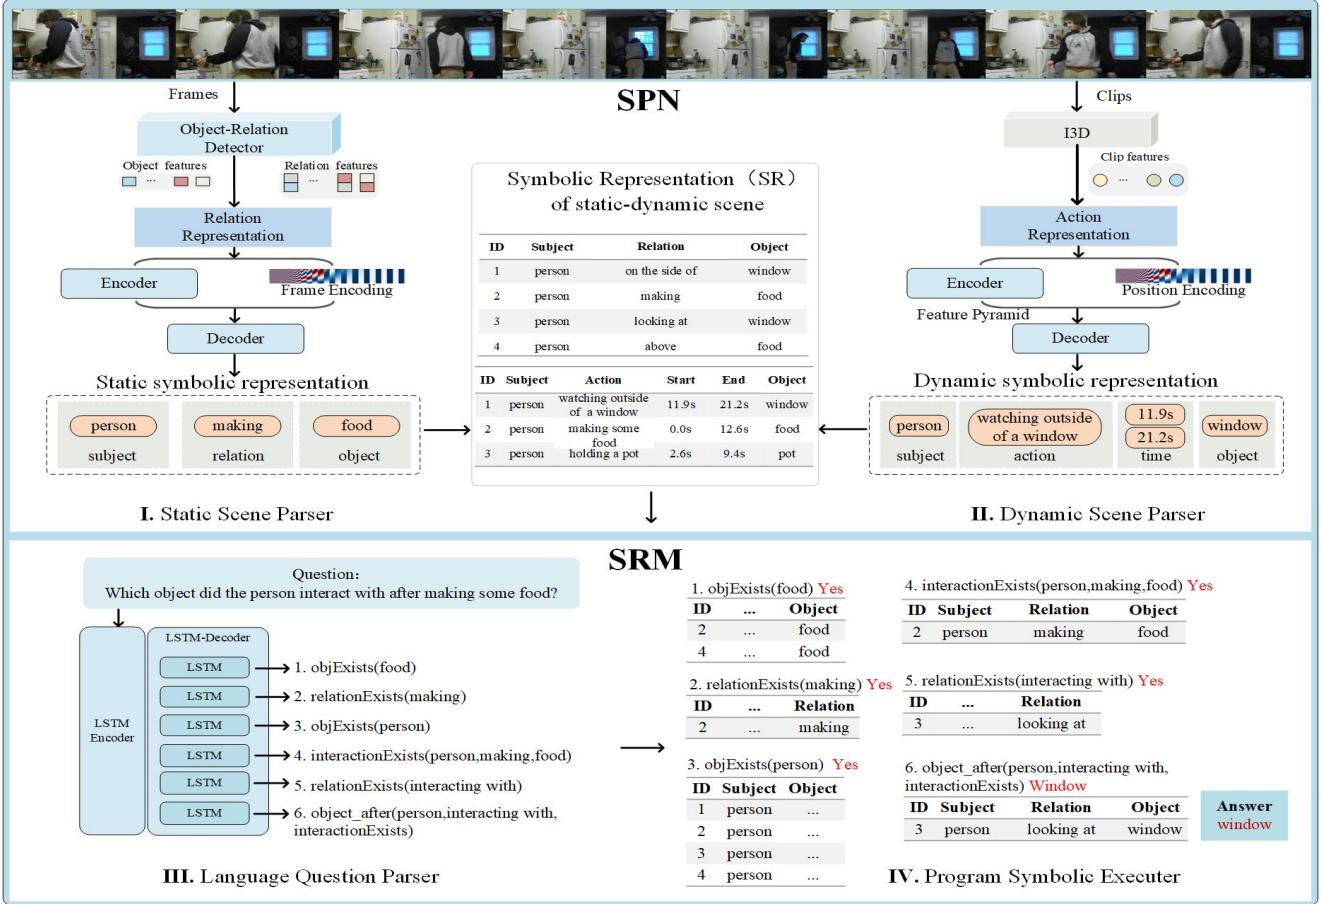


Figure 2. NS-VideoQA uses SPN (I,II) to convert the input video into SR, then uses SRM (III,IV) to decompose the compositional question into a program, and applies reasoning rules iteratively on SR according to the program, finally generates the answer of the compositional question.

follows:

$$\begin{aligned} Z_j^{(1)} &= \text{Att}_{dec}(Q = K = Z_j + E_{pos}, V = Z_j) \\ Z_j^{(n)} &= \text{Att}_{dec}(Q = K = V = Z_j^{(n-1)}) \end{aligned} \quad (4)$$

Where $E_{pos} \in \mathbb{R}^{\eta \times 1936}$ is a temporal position encoding.

The output of decoder is represented as $Z_j^{(L)} = \{z_j^2, \dots, z_j^N\}$, we predict the relationship s_{rel} between o_t^1 and o_t^i from z_j^i . Finally, we obtain the symbolic representation of static scene about the given video v :

$$\begin{aligned} R_{static}(v) &= \varphi(v) = \bigcup_{t=1}^T \rho_t \\ \rho_t &= \{e_2, \dots, e_N\} \\ e_i &= (s_{subj}, s_{rel}, s_{obj}), i=2, \dots, N \end{aligned} \quad (5)$$

DYNAMIC SCENE PARSER ψ . Inspired by ActionFormer [35], we utilize the pretrained I3D [36] model to extract clip features $X = \{x_1, x_2, \dots, x_T\}$, where T is the number of clips.

Action representation. A projection function embeds each clip feature x_i into a D -dimensional space.

$$E = [\text{Conv}(x_1), \text{Conv}(x_2), \dots, \text{Conv}(x_T)]^T \quad (6)$$

Where $\text{Conv}(x_i) \in \mathbb{R}^D$ consists of $conv1d$ with kernel size = 3 and stride = 1 and ReLU.

Encoder. We utilize encoder E , to obtain the multi-scale action representation $Z = \{Z^{(1)}, \dots, Z^{(n)}\}$. The query Q , key K , and value V in the encoder share the same input, and the output of the n -layer encoder is:

$$\begin{aligned} Z^{(n)} &= \downarrow \left(\text{Att}_{enc} \left(Q = K = V = Z^{(n-1)} + E_{pos} \right) \right) \\ Z^{(0)} &= E \end{aligned} \quad (7)$$

Where $E_{pos} \in \mathbb{R}^D$ is a position embedding, $\downarrow(\cdot)$ is the down-sampling operator, with a strided depthwise $conv1d$. We use 2x downsampling in this work.

Decoder. The decoder is a lightweight convolutional network with classification and regression heads. It predicts the category, starting and ending time of actions $\{y_1, \dots, y_N\}$ from Z , where $y_i =$

$(s_{sbj}, s_{act}, s_{obj}, t_{start}, t_{end})$, $i = 1, \dots, N$. The symbolic representation of dynamic scene is represented as:

$$R_{dynamic}(v) = \psi(v) = \{y_1, \dots, y_N\} \quad (8)$$

To get the SR of static-dynamic scene, we combine $R_{dynamic}(v)$ and $R_{dynamic}(v)$ according to Eq.1.

3.3. Symbolic Reasoning Machine

SRM decomposes compositional questions with a top-down process and infers answers with a bottom-up process. To realize the inference of leaf (*ObjectExists*, etc.) and compositional questions (*Choose*, etc.), SRM takes SR and questions as input, and applies symbolic reasoning rules to the SR to obtain the final answer.

Language Question Parser. Same as in [12]. To ensure the completeness of NS-VideoQA, we briefly describe the question parser. To obtain program sequence $p = \{p_1, p_2, \dots, p_M\}$, the question parser utilizes BiLSTM [37] to encode a question $q = \{q_1, q_2, \dots, q_T\}$, where M denotes the length of the program, T denotes the length of the question. The middle representation of the encoded question q at time i as:

$$\begin{aligned} e_i &= \left[e_i^F, e_i^B \right], \text{ where} \\ e_i^F, h_i^F &= \text{LSTM}(\Phi_E(q_i), h_{i-1}^F), \\ e_i^B, h_i^B &= \text{LSTM}(\Phi_E(q_i), h_{i+1}^B) \end{aligned} \quad (9)$$

Where Φ_E is the encoder word embedding. $(e_i^F, h_i^F), (e_i^B, h_i^B)$ are the outputs and hidden vectors of the forward and backward networks.

The decoder is a network of LSTM+Attention [38]. The decoder generates a token p_t according to the previous token of output sequence y_{t-1} , then p_t and e_i are integrated into Attention to obtain a context vector c_t , which is a weighted sum of the hidden state:

$$\begin{aligned} p_t &= \text{LSTM}(\Phi_D(y_{t-1})), \alpha_{ti} \propto \exp(p_t^T W_A e_i), \\ c_t &= \sum \alpha_{ti} e_i \end{aligned} \quad (10)$$

Where Φ_D is the decoder word embedding. W_A is the attention weight matrix. To obtain the distribution of predicted program $p_t \sim \text{softmax}(W_O[y_t, c_t])$, the context vector and decoder output are integrated into a fully connected layer with softmax activation.

Program Symbolic Executor. Considering the polymorphism of compositional questions (where the reasoning rules corresponding to the same combination type can vary depending on the category of the sub-question). For a program p with n direct sub-questions $C_p = \{c_1, \dots, c_n\}$, assuming τ_p is the question type of p , we define $token_p := \{\tau_p, \tau_{c_1}, \dots, \tau_{c_n}\}$. Each token corresponds to a reasoning rule $rule_{token}$, with more details in the Supplementary.

The reasoning is conducted by applying $rule_{token_p}$ iteratively. At each step, $rule_{token_p}$ is applied to different

inputs determined by the question structure. For leaf questions, the input is $SR(v)$; for compositional questions, the input is the intermediate results from all descendant questions of q . The result of a question denoted as $a_q := \{ans_q, R_q\}$, packs two components: a symbolic answer $ans_q \in (\{yes, no\} \cup V_{obj} \cup V_{rel} \cup V_{action})$, and symbolic representations $R_q \in SR(v)$ filtered for the question. All intermediate results obtained previously are recorded in a trace T , implemented as a dictionary. Thus, the reasoning step for question q becomes:

$$a_q = \begin{cases} \text{rule}_{token_p}(SR(v)), & q \in Q_l \\ \text{rule}_{token_p}(T), & q \in Q_c \end{cases} \quad (11)$$

Where Q_l is leaf question, Q_c is compositional question.

We design a program execution algorithm with dynamic programming, as Algorithm 1. For compositional questions, first, the trace is updated iteratively for each sub-question; then, the chosen rule is applied to the trace; finally, the obtained intermediate result is appended to the trace.

Algorithm 1 Program Execution

```

Input: program  $p$ , scene representation  $SR(v)$ , trace  $T_{in}$ 
Output: trace  $T$ 
# initialization
 $token \leftarrow \text{get\_token}(p)$ 
 $rule \leftarrow \text{get\_rule}(token)$ 
 $T \leftarrow T_{in}$ 
# reasoning
if  $p$  is leaf question then
     $a \leftarrow \text{rule}(SR(v))$ 
else
    for  $f$  in subprogram of  $p$  do
         $T \leftarrow T \cup \text{Program\_Execution}(f, SR(v), T)$ 
    end for
     $a \leftarrow \text{rule}(T)$ 
end if
# integration
 $T \leftarrow T \cup \{ (token, a) \}$ 
return  $T$ 

```

4. Experiments

To demonstrate the effectiveness of NS-VideoQA, our experimental design is as follows: First, we introduce the dataset and evaluation metrics. Secondly, we analyze the experimental results of the sub-questions, progressively presenting evidence to validate the reasoning capabilities of our model. Finally, we show the visualization process of model reasoning questions.

4.1. Dataset and Metrics

Dataset. The dataset of AGQA Decomp[14] is organized into 14 distinct sub-question categories, with each

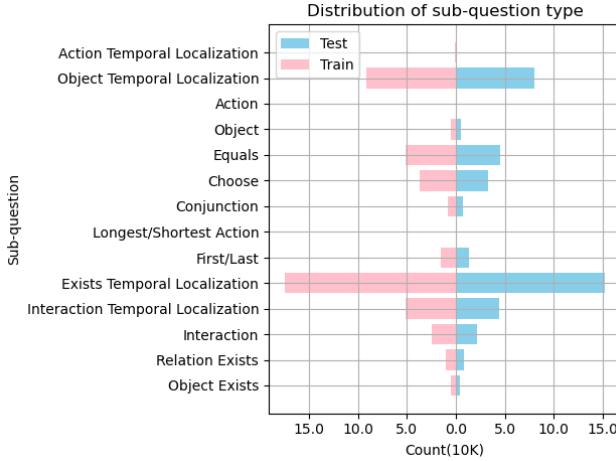


Figure 3. A chart displays the distribution of sub-question types in the train and test sets.

category linked to a specific function token. Leveraging indirect references of sub-questions through hand-designed reasoning rules enables the generation of corresponding questions and programs. The specific distribution of sub-question categories for training and testing is shown in Figure 3.

Metrics. For comprehensive analysis, we study the reasoning capability of our NS-VideoQA via measuring the accuracy, compositional accuracy (CA), right for the Wrong Reasons (RWR), delta, and internal consistency (IC) following AGQA Decomp[14]. CA measures the proficiency of the model in compositional reasoning, i.e., to test if the model can answer a question correctly based on correct reasoning results of its sub-questions. RWR scrutinizes the irrationality of reasoning performance by checking if the model guessed the answer correctly but based on wrong intermediate results from sub-questions. Delta represents the discrepancy between RWR and CA, and when RWR surpasses CA with a positive delta value, it indicates incorrect reasoning. IC quantifies the ability to avoid self-contradictions. A higher IC value suggests that the model exhibits compositional consistency satisfying logical reasoning rules. Precise definitions of these indicators and experimental details can be found in the Supplementary.

4.2. Performance on NS-VideoQA

Table 1 summarizes the comparative results across various sub-question categories. Here, Most-Likely [14] generates the most probable answer dependent on language biases. HME [39], HCRN [40], and PSAC [18], are purely neural VideoQA models. Oracle SR represents the reasoning performance of our SRM with ideal SR as inputs. Human is the annotators who agrees with ground-truth answers.

Our model achieves an average accuracy of 77.63% on

Table 1. We report accuracy values of sub-question type on AGQA Decomp. “-” indicates there were no results for a given type, bold numbers indicate the best accuracy.

Sub-question Type	Accuracy ↑						
	Most-Likely	HME	HCRN	PSAC	Ours	Oracle SR	Human
Object Exists	50.00	45.36	31.69	88.10	96.26	99.92	92.00
Relation Exists	50.00	51.33	61.72	-	100	72.94	99.00
Interaction	50.00	50.00	50.24	41.88	63.35	92.24	88.00
Interaction Temporal Loc.	50.00	50.00	50.85	39.37	50.76	73.80	96.00
Exists Temporal Loc.	50.00	50.00	45.27	39.55	-	-	92.00
Object Temporal Loc.	-	-	25.43	22.85	35.58	68.55	-
Action Temporal Loc.	-	-	4.06	3.70	67.17	92.03	-
Longest/Shortest Action	3.57	1.79	0.00	1.79	53.68	100	76.00
Action	-	-	0.00	16.67	57.14	100	-
Object	-	-	26.23	18.98	42.64	79.69	-
Choose	26.03	12.24	20.62	26.88	31.86	84.67	88.00
Equals	49.92	-	50.00	31.66	53.52	63.33	70.00
Conjunction	50.00	50.00	50.15	44.81	51.07	53.89	76.00
First/Last	3.79	2.93	0.00	4.31	13.19	46.01	88.00
Overall	3.31	8.28	12.43	14.98	23.51	74.52	84.36

leaf questions (*Object Exists* and *Relation Exists*, and *Interaction* categories). HCRN and HME achieve approximately 50% accuracy, while PSAC achieves an average accuracy of 76.66%. We conjecture that the accuracy disparity between purely neural models and neural-symbolic models remains minor in the absence of complex, multi-hop questions. To further investigate this idea, we provide accuracy results for each binary question type in the Supplementary. For the *Temporal Loc.* category, NS-VideoQA exhibits higher accuracy compared to HCRN, HME, and PSAC, especially in *Action Temporal Loc.*, *Longest/Shortest Action*, and *Action* categories. This can be attributed to the limited action detection and temporal localization capabilities of the above models.

For more complex categories, such as *Choose*, *Equals*, *Conjunction*, and *First/Last* categories, NS-VideoQA demonstrates superior performance compared to HCRN, HME, and PSAC by approximately 1% to 7%. The *Choose* type requires selecting objects or actions, while the *Equals* type involves evaluating the similarity between two objects or actions. The *Conjunction* type consists of And and Xor operations, and the *First/Last* type involves identifying the first/last object or action. Across these four types, NS-VideoQA consistently outperforms the other models, highlighting its strong capabilities in compositional spatio-temporal reasoning.

In the *Overall* results, NS-VideoQA achieves an accuracy of 23.51%, which is still far from human-level performance. However, it significantly improves the ability of compositional spatio-temporal reasoning compared to other models. To further analyze the reasoning capability of NS-VideoQA, we consider additional metrics, namely CA, RWR, and IC, as shown in Table 2. These metrics provide a more comprehensive understanding of NS-VideoQA’s performance and its effectiveness in capturing and reasoning about complex spatio-temporal information.

Table 2. We report compositional accuracy (CA), right for the wrong reasons (RWR), delta (RWR-CA),internal consistency (IC) values. The bold numbers indicate the best results, and the underlined numbers represent the results of the question categories that require attention.N/A indicates there were no valid compositions for a given type.

Sub-question Type	CA ↑				RWR ↓				Delta ↓				IC ↑			
	HME	HCRN	PSAC	Ours	HME	HCRN	PSAC	Ours	HME	HCRN	PSAC	Ours	HME	HCRN	PSAC	Ours
Object Exists	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
Relation Exists	N/A	92.86	100	85.39	0.00	52.46	N/A	N/A	N/A	-40.40	N/A	N/A	50.00	46.18	50.00	50.00
Interaction	N/A	67.31	71.47	87.13	19.58	50.61	63.91	29.62	N/A	-16.70	-7.57	-57.51	50.00	48.62	47.13	98.42
Interaction Temporal Loc.	100	53.69	53.00	38.20	50.08	49.82	24.22	52.20	-49.92	-3.87	-28.78	14.00	31.25	32.37	27.09	32.40
Exists Temporal Loc.	0.00	94.79	89.38	-	1.26	45.45	30.33	-	1.26	-49.34	-59.05	-	98.28	33.31	27.10	-
Object Temporal Loc.	60.00	37.55	35.10	80.46	24.75	29.36	23.91	15.08	-35.25	-8.19	-11.20	-63.58	77.04	56.71	39.70	33.65
Action Temporal Loc.	N/A	8.74	8.46	N/A	1.62	7.88	2.50	66.85	N/A	-0.86	-5.96	N/A	N/A	52.65	15.16	66.10
Longest/Shortest Action	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A	N/A
Action	N/A	0.00	28.57	57.14	0.00	0.00	N/A	N/A	N/A	0.00	N/A	N/A	14.29	14.92	14.29	15.38
Object	N/A	30.09	28.34	46.22	24.07	34.33	N/A	0.02	N/A	4.23	N/A	-46.20	14.29	14.67	14.29	24.24
Choose	N/A	39.12	42.96	39.07	43.28	44.75	48.57	31.26	N/A	5.62	5.61	-7.81	0.00	13.7	16.73	49.49
Equals	46.47	30.94	31.65	91.67	50.58	50.63	31.25	49.52	4.11	19.69	-0.40	-42.14	48.9	20.77	1.36	46.93
Conjunction	70.00	70.23	73.67	<u>63.33</u>	37.53	41.70	35.51	23.72	-32.47	-28.53	-38.17	-39.61	50.00	38.71	19.56	48.52
First/Last	47.01	0.03	70.66	81.66	10.02	0.01	7.79	3.64	-36.99	-0.02	-62.87	-78.02	36.80	0.02	75.56	26.65
Overall	23.10	37.52	45.23	51.44	18.77	29.07	19.14	23.00	-4.33	-8.45	-26.09	-28.44	35.67	26.62	24.85	37.04

Table 3. We report the performances of HME, HCRN, and PSAC on the RWR- n metrics (the lower the better), where n represents the number of incorrectly answered sub-questions for a composition, conditioned on the given parent question type. N/A occurs when none sub-question of this type has n sub-questions misanswered by that model.

Sub-question Type	HME					HCRN					PSAC					Ours				
	RWR-1	RWR-2	RWR-3	RWR-4	RWR-5	RWR-1	RWR-2	RWR-3	RWR-4	RWR-5	RWR-1	RWR-2	RWR-3	RWR-4	RWR-5	RWR-1	RWR-2	RWR-3	RWR-4	RWR-5
Object Exists	N/A	N/A	N/A	N/A																
Relation Exists	0.00	N/A	N/A	N/A	N/A	52.46	N/A	N/A	N/A	N/A										
Interaction	28.30	46.88	11.76	N/A	N/A	62.29	52.78	27.33	N/A	N/A	63.91	N/A	N/A	N/A	31.71	19.68	N/A	N/A	N/A	
Interaction Temporal Loc.	100	99.94	70.26	99.63	3.48	45.36	56.64	60.68	49.10	6.88	26.12	21.27	24.59	24.27	11.31	36.59	27.71	70.03	69.78	10.04
Exists Temporal Loc.	4.22	1.19	N/A	N/A	N/A	65.05	17.55	N/A	N/A	N/A	32.02	9.20	N/A	N/A	0.06	18.54	N/A	N/A	N/A	
Object Temporal Loc.	45.86	24.80	18.55	2.63	N/A	35.84	27.39	10.74	4.74	N/A	35.30	14.49	0.00	N/A	N/A	-	-	-	-	
Action Temporal Loc.	1.62	N/A	N/A	N/A	N/A	7.88	N/A	N/A	N/A	N/A	2.50	N/A	N/A	N/A	66.85	N/A	N/A	N/A	N/A	
Longest/Shortest Action	N/A	N/A	N/A	N/A																
Action	0.00	N/A	N/A	N/A	N/A	0.00	N/A	N/A	N/A	N/A										
Object	36.13	23.92	N/A	N/A	N/A	30.08	38.09	N/A	0.02	N/A	N/A	N/A	N/A							
Choose	44.35	18.94	N/A	N/A	N/A	44.80	44.03	N/A	N/A	N/A	55.60	32.46	N/A	N/A	29.26	15.62	N/A	N/A	N/A	
Equals	58.21	49.10	N/A	N/A	N/A	50.64	50.62	N/A	N/A	N/A	36.01	20.78	N/A	N/A	50.54	45.29	N/A	N/A	N/A	
Conjunction	46.81	24.35	N/A	N/A	N/A	49.27	24.74	N/A	N/A	N/A	51.58	6.35	N/A	N/A	44.89	79.46	N/A	N/A	N/A	
First/Last	10.02	N/A	N/A	N/A	N/A	0.01	N/A	N/A	N/A	N/A	7.79	N/A	N/A	N/A	3.64	N/A	N/A	N/A	N/A	
Overall	38.97	19.14	40.78	54.57	38.97	3.48	35.99	45.88	38.43	6.88	38.01	20.85	24.55	24.27	11.31	22.54	22.02	54.56	49.54	10.04

4.3. Performance on Choose, Equals, Conjunction, and First/Last

To further investigate the reasoning challenges of NS-VideoQA, we assess its ability to achieve high accuracy, and the probability of guessing the answer correctly based on incorrect reasoning. This evaluation is conducted using the CA, RWR, and Delta metrics, as presented in Table 2.

In the *Choose* type, NS-VideoQA exhibits a CA value of 39.07%, indicating that the model struggles with this category. Although its RWR value of 31.26% does not surpass HCRN and PSAC models, the Delta value of -7.81% suggests that NS-VideoQA is more likely to perform correct reasoning. Table 3 further validates this point, as NS-VideoQA exhibits lower RWR-1 and RWR-2 values compared to HCRN and PSAC. On the other hand, HCRN and PSAC have CA values higher than NS-VideoQA by 0.05% and 3.89%, respectively. However, their positive Delta values of 5.62% and 5.61% indicate that these models tend

to guess answers correctly based on incorrect reasoning. The values of CA, RWR, and Delta reveal the challenges faced by HME, HCRN, and PSAC in *Equals* type, with values of (46.47%, 30.94%, 31.56%), (50.58%, 50.63%, 31.25%), and (4.11%, 19.69%, -0.4%) respectively. In contrast, our model achieves 91.67%, 49.52%, and -42.14% respectively, indicating that NS-VideoQA has compositional spatio-temporal reasoning capability. Similar conclusions can be drawn for the *First/Last* type.

HME, HCRN, and PSAC demonstrate Delta values of -32.47%, -28.53%, and -38.17%, respectively, indicating their proficiency in correctly reasoning on *Conjunction* compared to *Choose* and *Equals*. NS-VideoQA exhibits a Delta value of -39.61%, suggesting that errors in intermediate reasoning have a minimal negative impact on its overall performance. This finding further highlights the robustness of NS-VideoQA based on reasoning rules and avoid relying on data biases. Additionally, we assess the model's self-

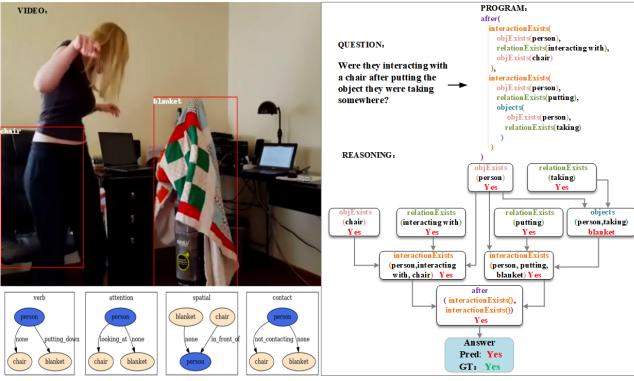


Figure 4. The reasoning for the *Interaction Temporal Loc.* type. **Top left:** the red boxes denote the detected objects “chair” and “blanket”. **Bottom left:** the static symbolic representation of verb, attention, spatial, and contact. **Right:** the answer “Yes” is obtained by reasoning based on the program.

consistency by evaluating the IC metrics in Table 2. To further prove the effectiveness of NS-VideoQA, we report the results of compositional types in Supplementary.

4.4. Visualization of execution traces and static-dynamic SR

The inherent visual interpretability of NS-VideoQA provides insights into SPN learning and SRM reasoning, leading to a detailed understanding of the mechanism inside. To illustrate, we visualize two instances of inference process as Figure 4 and Figure 5. The *Interaction Temporal Loc.* type shown in Figure 4 shows the reasoning involved with static SR. In the question, the phrase “they were taking” become a constraint to “the object”. With the relationship (person, taking, blanket) identified by SPN in previous frames, the SRM clarifies that “the object” refers to “blanket”. Thus, in the next reasoning step, the action description “putting the object” instantized to “putting the blanket”, enables the SRM to query the correct action instance from SR for the subsequent temporal reasoning.

Figure 5 shows the temporal reasoning appertain to dynamic SR. For two action instances specified by the given question, the dynamic SR represents their temporal groundings. By calculating and comparing their length, the SRM chooses the shorter one as the final answer in response to *Action Shorter* type. The above examples demonstrate: (i) The reasoning rules are conditioned by sub-question types and intermediate results. Thus, the polymorphism of the program executor is necessary. (ii) The reasoning process depends on both static and dynamic scene representations. Without any of them, the answer could not be deduced correctly. (iii) From the reasoning process, NS-VideoQA enjoys compositional spatio-temporal reasoning ability, interpretability, and generates human-readable execution traces.

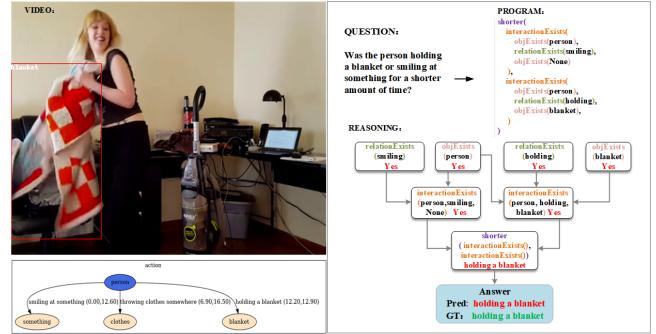


Figure 5. The reasoning for the *Shortest Action* type. **Top left:** the red box denotes the detected object “blanket”. **Bottom left:** the symbolic representation of dynamic scene. **Right:** the answer “hold a blanket” is obtained by reasoning based on the program.

5. Conclusion

In this paper, the proposed NS-VideoQA improves the capability of compositional spatio-temporal reasoning in real-world by transforming the video into symbolic representation and conducting iterative reasoning. The proposed model performed extensive experiments on AGQA Decomp dataset, the CA, RWR, Delta, and CA results show that NS-VideoQA exhibits internal consistency and significantly improves the accuracy of compositional spatio-temporal questions. The quantitative and qualitative analysis leads to the following conclusions: (i) In simple questions without complex spatio-temporal reasoning, neural-symbolic and purely neural methods show minimal differences. However, in complex scenarios requiring compositional spatio-temporal reasoning, purely neural VideoQA methods often rely on data biases for answer generation, lacking robust reasoning capabilities. (ii) Accurate static symbolic representations clarify relationships and filter objects, while accurate dynamic symbolic representations indicate action chronologies, aiding precise localization of behaviors and action sequencing.

The NS-VideoQA has certain limitations. When the symbolic representations are inaccurate (*e.g.*, low accuracy in object detection, relationship recognition, and action localization), the accuracy of question answering in NS-VideoQA can be affected. Additionally, even if the symbolic representations are accurate, the reasoning capabilities of NS-VideoQA can be hindered if the reasoning rules are ambiguous or prone to ambiguity. We hope our paper could motivate future research in obtaining accurate symbolic representations and generating reasoning rules with more unambiguity.

References

- [1] Khushboo Khurana and Umesh Deshpande. Video question-answering techniques, benchmark datasets and evaluation metrics leveraging video captioning: a comprehensive survey. *IEEE Access*, 9:43799–43823, 2021. [1](#)
- [2] Guanglu Sun, Lili Liang, Tianlin Li, Bo Yu, Meng Wu, and Bolun Zhang. Video question answering: a survey of models and datasets. *Mobile Networks and Applications*, pages 1–34, 2021. [1](#)
- [3] Jihyeon Lee, Wooyoung Kang, and Eun-Sol Kim. Dense but efficient videoqa for intricate compositional reasoning. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 1114–1123, 2023. [1](#)
- [4] Hsin-Ying Lee, Hung-Ting Su, Bing-Chen Tsai, Tsung-Han Wu, Jia-Fong Yeh, and Winston H Hsu. Learning fine-grained visual understanding for video question answering via decoupling spatial-temporal modeling. 2022.
- [5] Piyush Bagad, Makarand Tapaswi, and Cees GM Snoek. Test of time: Instilling video-language models with a sense of time. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2503–2516, 2023.
- [6] Aisha Urooj, Hilde Kuehne, Bo Wu, Kim Chheu, Walid Bousselham, Chuang Gan, Niels Lobo, and Mubarak Shah. Learning situation hyper-graphs for video question answering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14879–14889, 2023. [1](#)
- [7] Jinpeng Wang, Yixiao Ge, Rui Yan, Yuying Ge, Kevin Qinghong Lin, Satoshi Tsutsui, Xudong Lin, Guanyu Cai, Jianping Wu, Ying Shan, et al. All in one: Exploring unified video-language pre-training. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6598–6608, 2023. [1, 2](#)
- [8] Difei Gao, Luwei Zhou, Lei Ji, Linchao Zhu, Yi Yang, and Mike Zheng Shou. Mist: Multi-modal iterative spatial-temporal transformer for long-form video question answering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14773–14783, 2023. [2](#)
- [9] Junbin Xiao, Pan Zhou, Tat-Seng Chua, and Shuicheng Yan. Video graph transformer for video question answering. In *European Conference on Computer Vision*, pages 39–58. Springer, 2022. [1, 2](#)
- [10] Vaishak Belle. Symbolic logic meets machine learning: A brief survey in infinite domains. In *International conference on scalable uncertainty management*, pages 3–16. Springer, 2020. [2](#)
- [11] Dongran Yu, Bo Yang, Dayou Liu, Hui Wang, and Shirui Pan. A survey on neural-symbolic learning systems. *Neural Networks*, 2023. [2](#)
- [12] Kexin Yi, Jiajun Wu, Chuang Gan, Antonio Torralba, Pushmeet Kohli, and Josh Tenenbaum. Neural-symbolic vqa: Disentangling reasoning from vision and language understanding. *Advances in neural information processing systems*, 31, 2018. [2, 5](#)
- [13] Kexin Yi, Chuang Gan, Yunzhu Li, Pushmeet Kohli, Jiajun Wu, Antonio Torralba, and Joshua B Tenenbaum. Clevrer: Collision events for video representation and reasoning. In *International Conference on Learning Representations*, 2019. [2](#)
- [14] Mona Gandhi, Mustafa Omer Gul, Eva Prakash, Madeleine Grunde-McLaughlin, Ranjay Krishna, and Maneesh Agrawala. Measuring compositional consistency for video question answering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 5046–5055, 2022. [2, 3, 5, 6, 13](#)
- [15] Dejing Xu, Zhou Zhao, Jun Xiao, Fei Wu, Hanwang Zhang, Xiangnan He, and Yueling Zhuang. Video question answering via gradually refined attention over appearance and motion. In *Proceedings of the 25th ACM international conference on Multimedia*, pages 1645–1653, 2017. [2](#)
- [16] Zhou Zhao, Xinghua Jiang, Deng Cai, Jun Xiao, Xiaofei He, and Shiliang Pu. Multi-turn video question answering via multi-stream hierarchical attention context network. In *IJCAI*, volume 2018, page 27th, 2018. [2](#)
- [17] Thao Minh Le, Vuong Le, Svetha Venkatesh, and Truyen Tran. Learning to reason with relational video representation for question answering. *arXiv preprint arXiv:1907.04553*, 2, 2019. [2](#)
- [18] Xiangpeng Li, Jingkuan Song, Lianli Gao, Xianglong Liu, Wenbing Huang, Xiangnan He, and Chuang Gan. Beyond rnns: Positional self-attention with co-attention for video question answering. In *Proceedings of the AAAI conference on artificial intelligence*, volume 33, pages 8658–8665, 2019. [2, 6, 14](#)
- [19] Zekun Yang, Noa Garcia, Chenhui Chu, Mayu Otani, Yuta Nakashima, and Haruo Takemura. Bert representations for video question answering. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 1556–1565, 2020. [2](#)
- [20] Aisha Urooj, Amir Mazaheri, Mubarak Shah, et al. Mmfbert: Multimodal fusion transformer with bert encodings for visual question answering. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 4648–4660, 2020. [2](#)
- [21] Min Peng, Chongyang Wang, Yu Shi, and Xiang-Dong Zhou. Efficient end-to-end video question answering with pyramidal multimodal transformer. *arXiv preprint arXiv:2302.02136*, 2023. [2](#)
- [22] Jipeng Zhang, Jie Shao, Rui Cao, Lianli Gao, Xing Xu, and Heng Tao Shen. Action-centric relation transformer network for video question answering. *IEEE Transactions on Circuits and Systems for Video Technology*, 32(1):63–74, 2020. [2](#)
- [23] Justin Johnson, Bharath Hariharan, Laurens Van Der Maaten, Judy Hoffman, Li Fei-Fei, C Lawrence Zitnick, and Ross Girshick. Inferring and executing programs for visual reasoning. In *Proceedings of the IEEE international conference on computer vision*, pages 2989–2998, 2017. [2](#)
- [24] Jiayuan Mao, Chuang Gan, Pushmeet Kohli, Joshua B Tenenbaum, and Jiajun Wu. The neuro-symbolic concept learner: Interpreting scenes, words, and sentences from natural supervision. In *International Conference on Learning Representations*, 2018. [2](#)

- [25] Ramakrishna Vedantam, Karan Desai, Stefan Lee, Marcus Rohrbach, Dhruv Batra, and Devi Parikh. Probabilistic neural symbolic models for interpretable visual question answering. In *International Conference on Machine Learning*, pages 6428–6437. PMLR, 2019. [2](#)
- [26] Zhenfang Chen, Jiayuan Mao, Jiajun Wu, Kwan-Yee Kenneth Wong, Joshua B Tenenbaum, and Chuang Gan. Grounding physical concepts of objects and events through dynamic visual reasoning. In *International Conference on Learning Representations*, 2020. [2](#)
- [27] Zhenfang Chen, Kexin Yi, Antonio Torralba, Josh Tenenbaum, and Chuang Gan. Comphy: Compositional physical reasoning of objects and events from videos. In *International Conference on Learning Representations*, 2022. [2](#)
- [28] Jie Lei, Licheng Yu, Tamara Berg, and Mohit Bansal. Tqvqa+: Spatio-temporal grounding for video question answering. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8211–8225, 2020. [3](#)
- [29] Zhou Yu, Dejing Xu, Jun Yu, Ting Yu, Zhou Zhao, Yuet-ting Zhuang, and Dacheng Tao. Activitynet-qa: A dataset for understanding complex web videos via question answering. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 9127–9134, 2019. [3](#)
- [30] Madeleine Grunde-McLaughlin, Ranjay Krishna, and Maneesh Agrawala. Agqa: A benchmark for compositional spatio-temporal reasoning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 11287–11297, 2021. [3](#)
- [31] Madeleine Grunde-McLaughlin, Ranjay Krishna, and Maneesh Agrawala. Agqa 2.0: An updated benchmark for compositional spatio-temporal reasoning. *arXiv preprint arXiv:2204.06105*, 2022. [3](#)
- [32] Yuren Cong, Wentong Liao, Hanno Ackermann, Bodo Rosenhahn, and Michael Ying Yang. Spatial-temporal transformer for dynamic scene graph generation. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 16372–16382, 2021. [3](#)
- [33] Han Hu, Jiayuan Gu, Zheng Zhang, Jifeng Dai, and Yichen Wei. Relation networks for object detection. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3588–3597, 2018. [3, 11](#)
- [34] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017. [3](#)
- [35] Chen-Lin Zhang, Jianxin Wu, and Yin Li. Actionformer: Localizing moments of actions with transformers. In *European Conference on Computer Vision*, pages 492–510. Springer, 2022. [4, 12](#)
- [36] Joao Carreira and Andrew Zisserman. Quo vadis, action recognition? a new model and the kinetics dataset. In *proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 6299–6308, 2017. [4, 11](#)
- [37] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997. [5](#)
- [38] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*, 2014. [5](#)
- [39] Chenyou Fan, Xiaofan Zhang, Shu Zhang, Wensheng Wang, Chi Zhang, and Heng Huang. Heterogeneous memory enhanced multimodal attention model for video question answering. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 1999–2007, 2019. [6, 14](#)
- [40] Thao Minh Le, Vuong Le, Svetha Venkatesh, and Truyen Tran. Hierarchical conditional relation networks for video question answering. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 9972–9981, 2020. [6, 15](#)

Supplementary

The supplementary sections provide more detail on the methods and experiments described in our paper.

To help results reproduction, we report:

- Implementation details of NS-VideoQA in **A**,
- Feature extraction process in **B**,
- Training methods in **C**,
- Details of metrics in **D**.

Addition to the experimental results, we report:

- Performance on composition questions in **E**,
- Extra examples in **F**,
- Reasoning rules of AGQA Decomp in **G**.

A. Implementation Details

We report our implementation details including training settings, hyper-parameters, and input/output dimensions for each part of our NS-VideoQA. The implementation details for Object-Relation Detector, Faster RCNN, ActionFormer and Bi-LSTM Encoder-Decoder are listed in Table 4, Table 5, Table 6, and Table 7 respectively.

B. Feature Extraction

Static Scene Parser. Faster R-CNN based on ResNet101 is used as object detection backbone, and Relation Network [33] is applied to the instance recognition stage of Faster R-CNN. Each frame I_t is fed into Faster R-CNN to detect 100 objects, and obtain the 2048D object feature vectors, bounding boxes, and class distributions. The class distributions are linearly transformed into 200D semantic features by a linear matrix $W_s \in \mathbb{R}^{36 \times 200}$. The instance recognition stage computes relation features based on each object feature and bounding box.

For the 100 objects detected on the t -th frame $\{v_t^i, b_t^i\}_{i=1}^{100}$, the object weight $w_{t,v}^{ji}$ and box weight $w_{t,b}^{ji}$ of the i -th object with respect to any other object j are calculated as follows:

$$w_{t,v}^{ji} = \frac{\text{dot}(W_v v_t^j, W_q v_t^i)}{\sqrt{d_k}} \quad (12)$$

$$w_{t,b}^{ji} = \max \left\{ 0, W_b \cdot \varepsilon_b \left(b_t^j, b_t^i \right) \right\} \quad (13)$$

$$\varepsilon_b \left(b_t^j, b_t^i \right) = \left[\log \left(\frac{|x_j - x_i|}{w_j}, \frac{|y_j - y_i|}{h_j}, \frac{w_i}{w_j}, \frac{h_i}{h_j} \right) \right]^T \quad (14)$$

Where $\sqrt{d_k}$ is scaling factor. Finally, the relation feature r_t^i is computed as:

$$w_t^{ji} = \frac{w_{t,b}^{ji} \cdot \exp(w_{t,v}^{ji})}{\sum_k w_{t,b}^{ki} \cdot \exp(w_{t,v}^{ki})} \quad (15)$$

$$r_t^i = \sum_j w_t^{ji} \cdot (W_v v_t^j) \quad (16)$$

Table 4. Training settings, hyper-parameters, and input/output dimension of the Object-Relation Detector in Static Scene Parser.

Object-Relation Detector	key	value
training settings	learning rate	10^{-4}
	optimizer	AdamW
	epoch	10
	batch size	10 frames
	weight decay	10^{-3}
	dropout	0.1
hyper-parameters	encoder layer	1
	decoder layer	3
	window size η for transformer	2
input/output dimension	window stride for transformer	1
	object feature length	2048D
	relation feature length	2048D
	semantic feature length	200D
	feature length for transformer	1936D

Table 5. Training settings, hyper-parameters, and input/output dimension of the Faster RCNN in Static Scene Parser.

Faster RCNN	key	value
training settings	learning rate	10^{-4}
	optimizer	SGD
	epoch	20
	batch size	4 frames
	weight decay	10^{-4}
hyper-parameters	object numbers	100
	feature map length	1024D
input/output dimension	object feature length	2048D
	relation feature length	2048D

Table 6. Training settings, hyper-parameters, and input/output dimension of the ActionFormer in Dynamic Scene Parser.

ActionFormer	key	value
training settings	learning rate	10^{-4}
	optimizer	Adam
	epoch	50
	warmup epoch	5
	batch size	2 clips
	weight decay	0.05
	loss coefficient λ	1
hyper-parameters	dropout	0.1
	feature pyramid layer	7
input/output dimension	downsampling ratio	2
	motion feature length	1024D

The parameters of static scene parser including RPN are fixed when training scene graph generation models.

Dynamic Scene Parser. The motion features are obtained by passing RGB clips through an I3D [36] network pre-trained on Kinetics400. Following the experimental

Table 7. Training settings, hyper-parameters, and input/output dimension of the Bi-LSTM Encoder-Decoder in Symbolic Reasoning Machine.

Bi-LSTM Encoder-Decoder	key	value
training settings	learning rate	7×10^{-4}
	optimizer	Adam
	iteration	21721
	batch size	64 questions
hyper-parameters	weight decay	0.9
	hidden layer	2
	hidden layer dimension	256D
	encoder layer	1
input/output dimension	decoder layer	1
	GLoVe word embedding	300D
input/output dimension	length of question sequence	49D
	length of program sequence	81D

setup in ActionFormer [35], we first resize the frames to scale the shortest side to 256; then crop to 224×224 at center. The clips are obtained by a sliding window with size 16 and stride 4. For each clip, we extract the activation value of the group 5 in the I3D network and apply a global average pooling to obtain a feature vector of 1024D.

C. Training NS-VideoQA

Training SPN. The Static Scene Parser and the Dynamic Scene Parser in the SPN can be trained separately, such modular architecture makes SPN compatible to the improvements of relation recognition and temporal action detection methods in future.

In the Static Scene Parser, a cross-entropy loss L_{obj} is applied to train object detection, and a multi-label margin loss L_r is introduced to optimize the detection of multiple relation for the same subject-object pair. For a subject-object pair p , let R^+ and R^- be the relation set exists and not exists between p in ground-truth, let $\alpha(p, r)$ be the predicted probability that the subject-object pair p has relation r , then the L_r is calculated as:

$$L_r(p, R^+, R^-) = \sum_{r \in R^+} \sum_{r' \in R^-} \max(0, 1 - \alpha(p, r) + \alpha(p, r')) \quad (17)$$

The total loss L_s of the Static Scene Parser is:

$$L_s = L_r + L_{obj} \quad (18)$$

In the Dynamic Scene Parser, for each action, its category and start/end time are predicted by classification and regression, respectively. Thus, the classification loss L_{cls} for the action category is a Focal loss, and the regression loss L_{reg} for the action start/end time is a DIoU loss.

The total loss L_d of the Dynamic Scene Parser is:

$$L_d = \sum_L \sum_t (L_{cls} + \lambda(0, 1)L_{reg}) / T^+ \quad (19)$$

Where L refers to the number of layers in the feature pyramid, T^+ is the number of positive samples, $\lambda = 1$ is a coefficient to balance the classification and regression.

Training SRM. In SRM, a Bi-LSTM Encoder-Decoder is trained as the Language Question Parser to generate the programs from questions, where each element of the output sequence refers to a program statement corresponding to a reasoning step. We divide all program statements into a total of $J = 1210$ categories according to their question type and sub-question types. Treating the prediction of each element as a classification problem, the loss function of the Language Question Parser is an element-by-element cross-entropy loss:

$$L_{prog} = - \sum_{j=1}^J \mathbb{I}\{y_p = j\} \log(p_j) \quad (20)$$

Where, p_j is the predicted probability that the current element belongs to the j -th category, y_p is the ground-truth label of the current element.

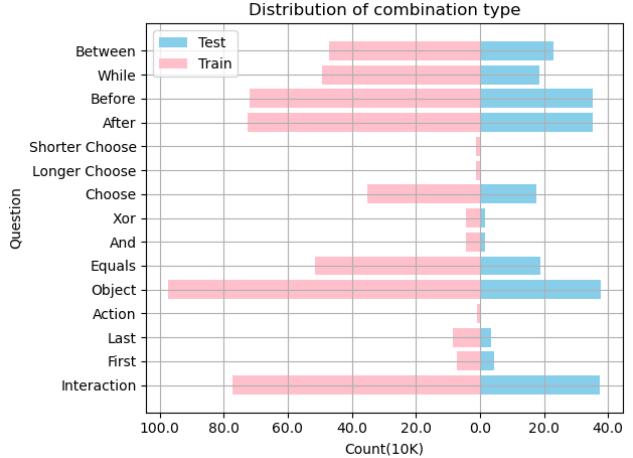


Figure 6. A chart displays the distribution of combination types in the train and test sets.

D. Metrics

The performances were evaluated using the metrics CA, RWR, Delta, and IC, and Accuracy proposed in the AGQA-Decomp benchmark. The definitions of these metrics are described above.

Notations. Here, we denote each test sample in the test set S as (v, q) , where v and q refer to video and question respectively. Let $gt(v, q)$ be the ground-truth answer of (v, q) , and $f(v, q)$ be the answer predicted by model f . For convenience, we define the video-question pairs correctly answered by f in S as $a(S, f)$, formally, $a(S, f) = \{(v, q) | (v, q) \in S, f(v, q) = gt(v, q)\}$. Additionally, for compositionality analysis, we denote the set of compositional questions as Q_c , where each $q \in Q_c$ has a set of sub-questions $C(q)$.

Table 8. We report compositional accuracy (CA), right for the wrong reasons (RWR), delta (RWR-CA), internal consistency (IC) values respect to composition rules for HME, HCRN, PSAC and Ours model.

Composition Type	CA \uparrow				RWR \downarrow				Delta \downarrow				IC \uparrow			
	HME	HCRN	PSAC	Ours	HME	HCRN	PSAC	Ours	HME	HCRN	PSAC	Ours	HME	HCRN	PSAC	Ours
Interaction	100	56.12	56.07	77.43	36.71	52.65	41.75	50.75	-63.29	-3.47	-14.31	-26.68	50.00	48.61	39.69	61.14
First	48.16	0.02	72.66	92.63	8.62	0.01	7.90	1.72	-39.54	-0.01	-64.77	-90.91	37.90	0.03	74.96	52.12
Last	45.61	0.04	68.19	70.89	11.93	0.01	7.65	6.36	-33.68	-0.04	-60.55	-64.54	35.70	0.02	76.15	22.47
Action	N/A	8.61	8.88	57.14	1.59	7.69	2.50	26.95	N/A	-0.92	-6.38	-30.19	42.37	34.12	25.79	42.08
Object	N/A	33.63	31.90	31.64	24.73	32.47	27.56	15.04	N/A	-1.16	-4.34	-16.6	58.31	48.73	36.97	50.40
Equals	46.47	30.94	31.65	91.67	50.58	50.63	31.25	49.52	4.11	19.69	-0.40	-42.14	48.90	20.77	1.36	46.93
And	100	89.16	96.81	54.42	41.28	27.67	14.08	48.77	-58.72	-61.49	-82.73	-5.65	50.00	61.78	39.13	51.02
Xor	100	38.40	0.00	94.81	33.26	53.04	50.79	40.60	-66.74	14.63	50.79	-54.21	50.00	15.64	0.00	82.18
Choose	N/A	40.61	80.95	86.85	44.36	45.08	48.57	31.26	N/A	4.47	-32.38	-55.59	0.00	13.70	16.73	49.19
Longer Choose	N/A	33.68	42.65	0.00	19.19	33.78	N/A	N/A	N/A	0.10	N/A	N/A	N/A	N/A	N/A	N/A
Shorter Choose	N/A	33.96	43.08	0.00	18.26	32.34	N/A	N/A	N/A	-1.62	N/A	N/A	N/A	N/A	N/A	N/A
After	72.73	76.59	72.05	23.80	21.86	49.53	23.89	21.87	-50.87	-27.06	-48.15	-1.93	68.87	52.11	38.46	38.67
Before	63.16	77.50	71.70	21.69	21.27	47.66	23.69	20.49	-41.89	-29.84	-48.01	-1.20	68.84	51.89	35.76	48.30
While	75.00	89.99	83.03	7.80	8.43	47.20	26.42	7.49	-66.57	-42.79	-56.61	-0.31	59.45	54.9	46.24	27.92
Between	95.83	92.91	98.18	94.28	6.60	18.73	16.93	15.95	-89.23	-74.18	-81.25	-78.32	58.50	68.20	42.57	29.16
Overall	86.52	62.25	62.63	36.40	26.75	43.13	32.15	27.41	-59.76	-19.13	-30.48	-8.09	48.92	40.34	34.84	42.53

Table 9. We report the performances of HME, HCRN, and PSAC on the RWR-n metrics (the lower the better), where n represents the number of incorrectly answered sub-questions for a composition, conditioned on the given composition type.

Composition Type	HME			HCRN			PSAC			Ours		
	RWR-1	RWR-2	RWR-3									
Interaction	60.96	89.01	9.84	59.02	58.15	25.37	45.50	28.86	29.23	39.11	20.36	11.79
First	8.62	N/A	N/A	0.01	N/A	N/A	7.90	N/A	N/A	1.72	N/A	N/A
Last	11.93	N/A	N/A	0.01	N/A	N/A	7.65	N/A	N/A	6.36	N/A	N/A
Action	1.59	N/A	N/A	7.69	N/A	N/A	2.50	N/A	N/A	66.85	N/A	N/A
Object	40.33	22.12	N/A	33.06	31.76	N/A	30.46	24.53	N/A	8.89	16.65	N/A
Equals	58.21	49.10	N/A	50.64	50.62	N/A	36.01	20.78	N/A	50.54	45.29	N/A
And	100	0.00	N/A	32.99	17.22	N/A	18.70	5.07	N/A	40.94	29.01	N/A
Xor	14.98	100	N/A	61.47	31.93	N/A	76.07	7.19	N/A	16.88	98.03	N/A
Choose	44.35	81.48	N/A	45.08	45.10	N/A	55.60	32.46	N/A	29.26	55.62	N/A
Longer Choose	50.00	19.17	N/A	34.18	32.53	N/A						
Shorter Choose	0.00	18.27	N/A	31.35	35.26	N/A						
After	60.95	19.80	N/A	61.40	31.61	N/A	24.92	10.11	N/A	22.69	17.50	N/A
Before	60.49	19.18	N/A	59.64	29.78	N/A	24.98	7.58	N/A	23.30	21.03	N/A
While	64.97	6.63	N/A	64.41	19.14	N/A	27.06	8.43	N/A	9.48	6.39	N/A
Between	45.79	1.31	N/A	36.44	6.36	N/A	35.54	4.68	N/A	7.63	17.07	N/A
Overall	45.42	22.96	9.84	49.05	35.55	25.37	36.24	20.32	29.23	26.45	20.14	11.79

Compositional accuracy (CA). CA tests if the model can answer a question correctly based on correct intermediate results. Let S_{CA} be the set of video-question pairs that all direct sub-questions have answered correctly by f , i.e., $S_{CA} = \{(v, q) | (v, q) \in S, q \in Q_c, |a(C(q), f)| = |C(q)|\}$. The CA is calculated by:

$$CA(f) = \frac{|a(S_{CA}, f)|}{|S_{CA}|} \quad (21)$$

Right for the wrong reasons (RWR). RWR checks if the model guessed the answer correctly but based on wrong intermediate results. Let S_{RWR} be the set of video-question pairs that has some direct sub-question answered incor-

rectly by f , i.e., $S_{RWR} = Q_c - S_{CA}$. The RWR is calculated by:

$$RWR(f) = \frac{|a(S_{RWR}, f)|}{|S_{RWR}|} \quad (22)$$

In addition, following [14], we measure the detailed RWR of our NS-VideoQA by the RWR- n metric, where n represents the exact number of incorrectly answered sub-questions for a composition, i.e., to replace S_{RWR} with $S_{RWR-n} = \{(v, q) | (v, q) \in S, q \in Q_c, |a(C(q), f)| = |C(q)| - n\}$.

Delta. Delta represents the discrepancy between RWR and CA. The Delta is calculated by:

$$\Delta(f) = RWR(f) - CA(f) \quad (23)$$

Table 10. We present internal consistency (IC) scores for individual logical consistency rules for HCRN, HME, PSAC and the Ours model. The N/A indicates that the amount of data for a combination type is 0.

Consistency Check	Parent Answer	IC ↑			
		HME	HCRN	PSAC	Ours
Interaction	Yes	N/A	50.62	79.39	24.17
	No	100	46.6	N/A	98.11
First	Object	38.23	0.02	74.95	33.64
	Action	11.83	1.08	76.52	70.61
Last	Object	35.81	0.01	76.11	19.18
	Action	7.58	0.76	87.12	25.76
Action	After	98.15	57.41	21.37	44.30
	Before	98.42	47.89	8.95	87.89
Object	After	77.35	55.55	42.25	35.55
	Before	76.96	56.94	43.26	33.44
	While	78.35	58.99	43.87	35.01
	Between	75.49	55.37	29.42	30.59
Equals	Yes	0.00	5.18	2.72	5.45
	No	97.80	36.37	0.00	88.41
And	Yes	N/A	73.58	78.26	31.80
	No	100	49.99	N/A	70.23
Xor	Yes	N/A	8.95	0.00	66.51
	No	100	22.33	0.00	97.85
Choose	Temporal	0.00	0.47	0.00	0.86
	Object	0.00	11.54	0.00	96.52
	Action	0.00	29.11	N/A	50.19
After	Yes	N/A	48.78	2.08	76.29
	No	99.97	46.70	56.52	33.26
Before	Yes	N/A	50.64	1.98	78.37
	No	99.97	52.07	56.89	34.01
While	Yes	N/A	52.28	2.30	77.71
	No	100	53.43	70.18	13.52
Between	Yes	N/A	81.10	88.13	38.65
	No	100	68.12	99.73	18.25
Object	Yes	0.00	35.73	100	100
	No	100	66.95	0	69.69
Overall	-	87.32	44.54	59.49	42.81

Table 11. We report accuracy per ground-truth answer for each binary sub-question type expecting “Yes” or “No” answers for HCRN, HME, PSAC and Ours model. “-” indicates there were no results for a given type, N/A indicates that the amount of data for a question type is 0. HME particularly is biased towards “No” for all question types, while PSAC particularly is biased towards “Yes” for all question types.

Sub-question Type	Ground Truth	ACC ↑			
		HME	HCRN	PSAC	Ours
Object Exists	Yes	0.00	31.69	88.10	96.26
	No	N/A	N/A	N/A	N/A
Relation Exists	Yes	0.00	61.72	100	72.94
	No	N/A	N/A	N/A	N/A
Interaction	Yes	0.00	55.11	83.75	57.20
	No	100	45.37	0.00	86.30
Exists Temporal Loc.	Yes	0.00	64.01	79.10	-
	No	100	26.54	0.00	-
Interaction Temporal Loc.	Yes	0.00	28.12	78.73	21.64
	No	100	73.58	0.01	79.99
Conjunction	Yes	0.00	24.34	89.61	96.28
	No	100	75.96	0.00	5.85
Equals	Yes	0.00	61.64	62.00	10.96
	No	100	38.72	0.00	95.12

A positive Delta value indicates that for the model f , answering sub-questions incorrectly leads to higher accuracy, which reveals the irrational reasoning steps taken by the model f .

Internal Consistency (IC). IC quantifies the ability to avoid self-contradictions. To measure IC, a set of logical consistency rules Φ is given in AGQA-Decomp. For each rule $\phi \in \Phi$, each question q derives $R_\phi(q)$, a set of questions such that for $\forall q' \in R_\phi(q)$, the answer to q' can be inferred from the answer to q according to the rule ϕ . Then, the IC can be calculated by:

$$IC(f) = \frac{\sum_{\phi \in \Phi} IC_\phi(f)}{|\Phi|} \quad (24)$$

$$IC_\phi(f) = \frac{\sum_{(v,q) \in S, q \in Q_c} |a(R_\phi(q), f)|}{\sum_{(v,q) \in S, q \in Q_c} |R_\phi(q)|} \quad (25)$$

Accuracy. The Accuracy scores are calculated for each question type, and the *Overall* accuracy is the average accuracy of all question types. Formally, let T be the set of all question types. Each question type $t \in T$ has an answer set A_t . The questions of type t with a specific ground truth answer $g \in A_t$ are denoted as $Q_{t,g} \subseteq S$. Then, the accuracy score $Accuracy(f, t)$ for model f and question type t is calculated by:

$$Accuracy(f, t) = \frac{1}{|A_t|} \sum_{g \in A_t} \left(\frac{1}{|Q_{t,g}|} a(Q_{t,g}, f) \right) \quad (26)$$

The overall accuracy is:

$$Accuracy_{all}(f) = \frac{1}{|T|} \sum_{t \in T} Accuracy(f, t) \quad (27)$$

E. Performance on composition questions

In the experimental section sec. 4, we reported the evaluation results of Accuracy, CA, RWR, Delta, and IC on different sub-questions. The experimental results demonstrate that the NS-VideoQA model possesses a certain level of compositional reasoning ability and meets the internal consistency check. To further validate the compositional reasoning capability of NS-VideoQA, we conducted supplemental experiments on composite questions of AGQA Decomp. The distribution of combination types is illustrated in Figure 6.

Based on the experimental results presented in Tables 8, 9, 10, 11, it can be observed that the HME [39] model is biased towards “No” answers, while the PSAC [18] model is biased towards “Yes” answer. Therefore, these models exhibit limitations in their ability to reason over given compositions, often resulting in correct answers for the wrong reasons, which can be attributed to self-contradiction. To illustrate, let’s consider the *Interaction* composition. The HME model achieves the following values in Table 8: CA (100%), RWR (36.71%), Delta (-63.29%), and IC (50%).

In Table 9, the model achieves RWR-1 (60.96%), RWR-2 (89.01%), and RWR-3 (9.84%). Moving on to Table 10, we find the IC values for *Interaction*-Yes and *Interaction*-No as (N/A, 100%). Here, N/A indicates that there are no instances where *Interaction* compositions result in “Yes” answers. Furthermore, Table 11 shows that the Accuracy values for *Interaction*-Yes and *Interaction*-No are (0.00%, 100%). This is sufficient to prove the above conclusion.

Therefore, let’s briefly compare the differences between HCRN [40] and NS-VideoQA in terms of experimental results. Our model demonstrates a stronger advantage in compositional reasoning tasks. For the *First*, *Last*, *Action*, *Object*, and *Choose* compositions, our model surpasses HCRN in terms of CA values by 21.31%, 92.61%, 70.85%, -1.99%, and 46.24% respectively. Similarly, our model outperforms HCRN in IC values by 12.53%, 52.09%, 22.45%, 1.67%, and 35.79% respectively. Regarding the *And* and *Xor* compositions, our model excels in answering *Xor*, while HCRN performs better in answering *And*. In terms of *After*, *Before*, *While*, and *Between* compositions, the overall differences between our model and HCRN, as observed from Tables 8, 9, 10, 11, are not significant. Although HCRN achieves slightly higher CA and IC values compared to our model, our model exhibits relatively lower RWR values, with reductions of 27.66%, 27.17%, 39.71%, and 2.78% respectively. Additionally, our model shows significantly lower RWR-1 and RWR-2 values compared to HCRN, indicating that HCRN relies more on guessing the answers to compositional questions based on incorrect sub-question results. For *Overall*, our model surpasses HCRN by 2.19% in terms of IC values.

F. Extra Examples from AGQA Decomp

In this section, we will showcase the spatio-temporal reasoning process for a variety of question types in Figure 7, 8, 9, 10, 11, 12, 13, 14, such as *Interaction Temporal Loc.*, *Conjunction*, *First/Last* and *Equals* categories. These examples will include instances of correct reasoning and instances of erroneous reasoning in NS-VideoQA. In the incorrect examples, the differences in symbolic reasoning programs stem from the variations in the descriptions of the input natural language questions. Our Language Question Parser, which utilizes a Bi-LSTM Encoder-Decoder, remains uninterpretable internally. Fortunately, in the NS-VideoQA framework, the neural-symbolic integration architecture allows us to conduct error analysis around the reasoning process and pinpoint the issue to the Language Question Parser, inspiring future improvements in program generation methods.

G. Reasoning rules of SRM

In this section, we present the definitions of reasoning rules for each sub-question type, including functional descriptions and input/output specifications. Specifically, Tables 12, 13, 14 list the reasoning rules corresponding to all sub-questions, while Table 15 provides the definitions of input/output data types for the reasoning rules.

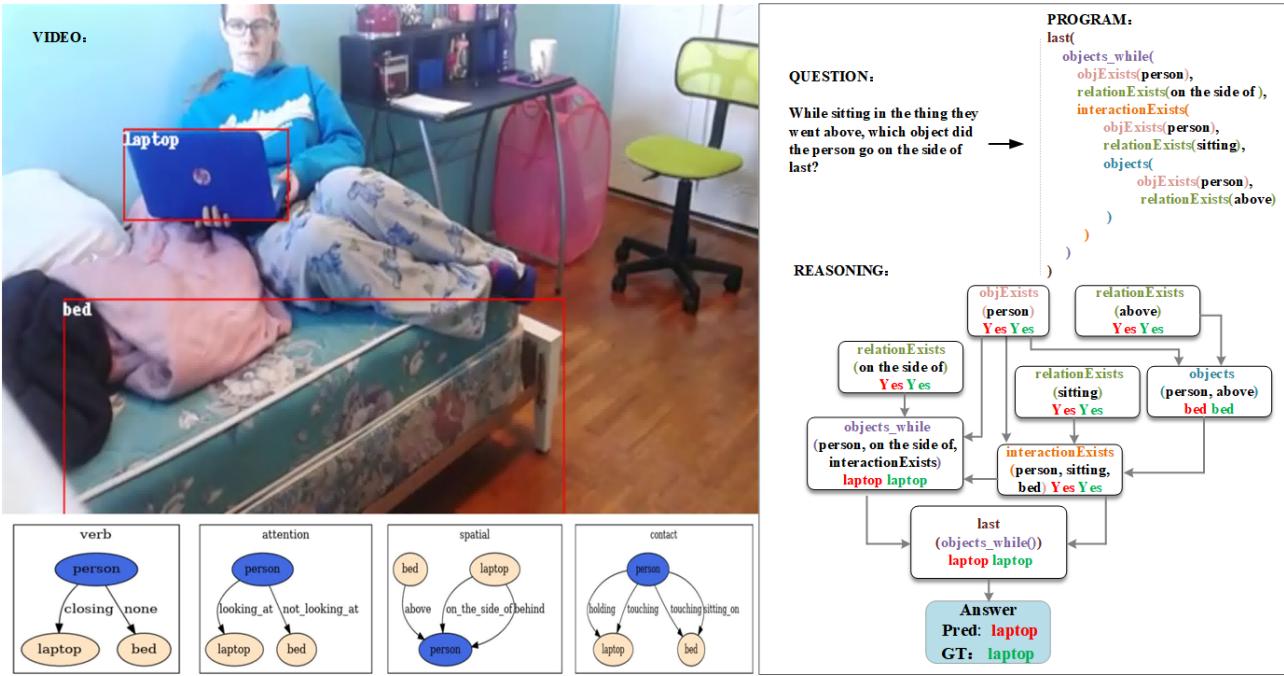


Figure 7. An example of a type *Last* inference process, where red indicates the prediction from NS-VideoQA and green represents the Ground Truth (GT).

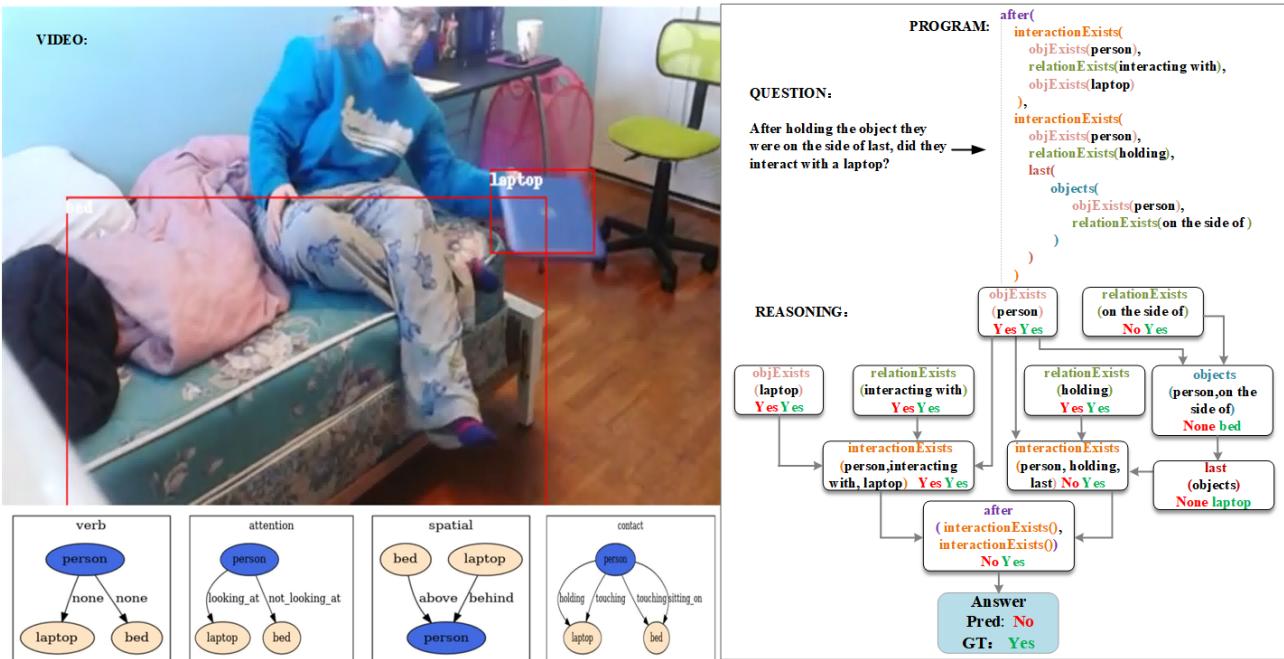


Figure 8. An example of a type *Interaction Temporal Loc.* inference process, where red indicates the prediction from NS-VideoQA and green represents the Ground Truth (GT).

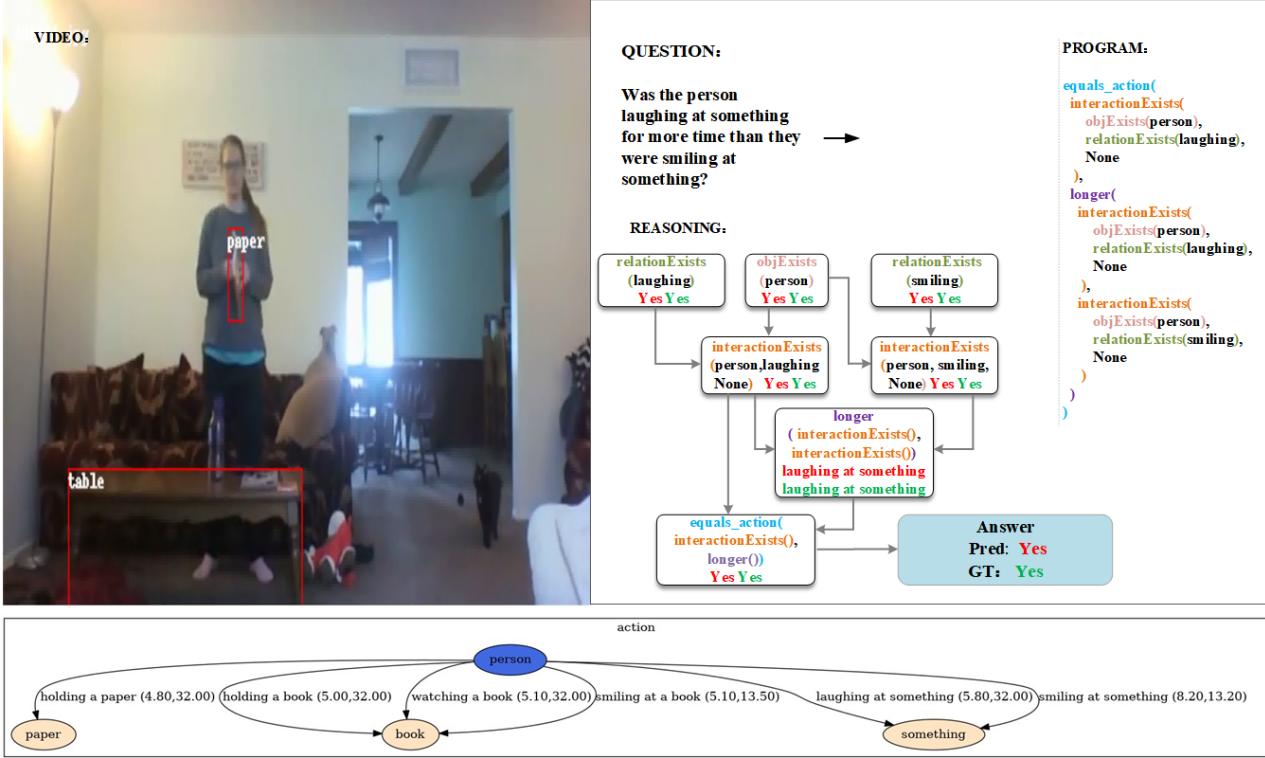


Figure 9. An example of a type *Equals* inference process, where red indicates the prediction from NS-VideoQA and green represents the Ground Truth (GT).

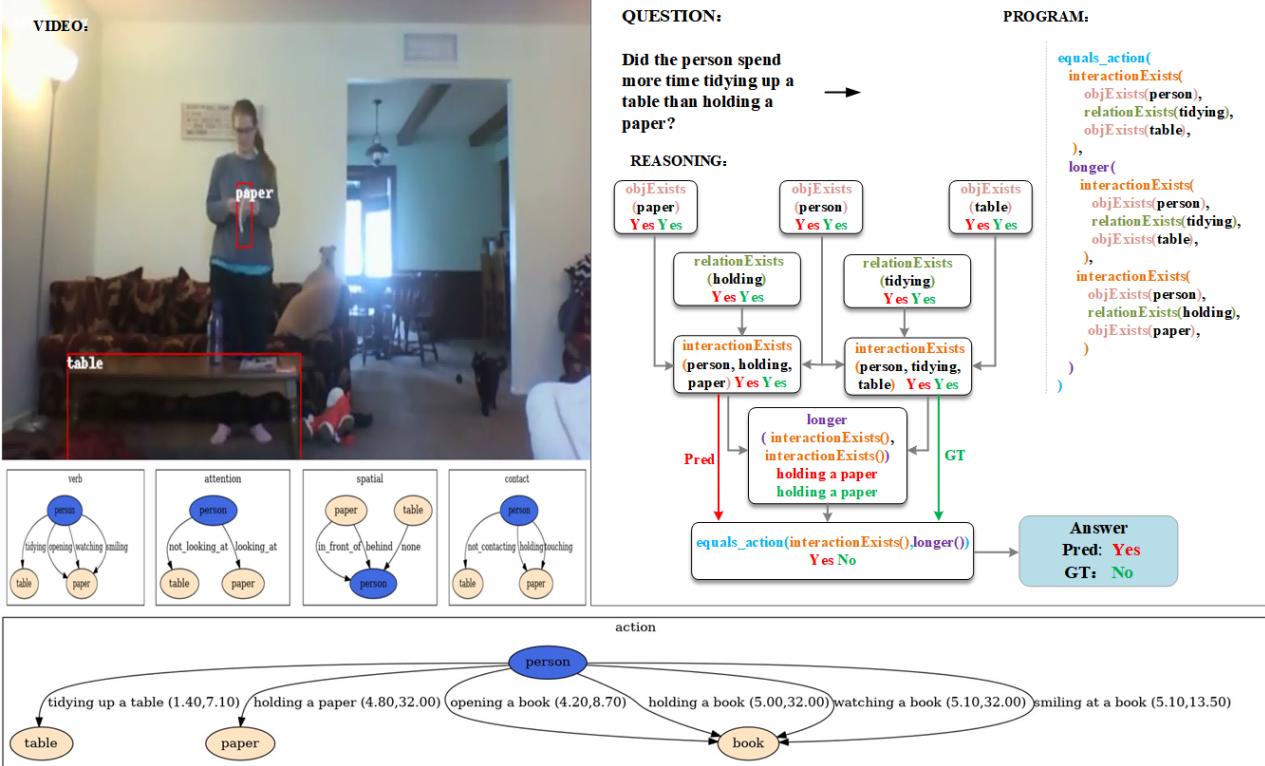


Figure 10. An example of a type *Equals* inference process, where red indicates the prediction from NS-VideoQA and green represents the Ground Truth (GT). The generation order of the two *interactionExists* sub-questions is swapped, resulting in passing incorrect parameters to the equals function (as indicated by the red arrow), leading to the opposite result.

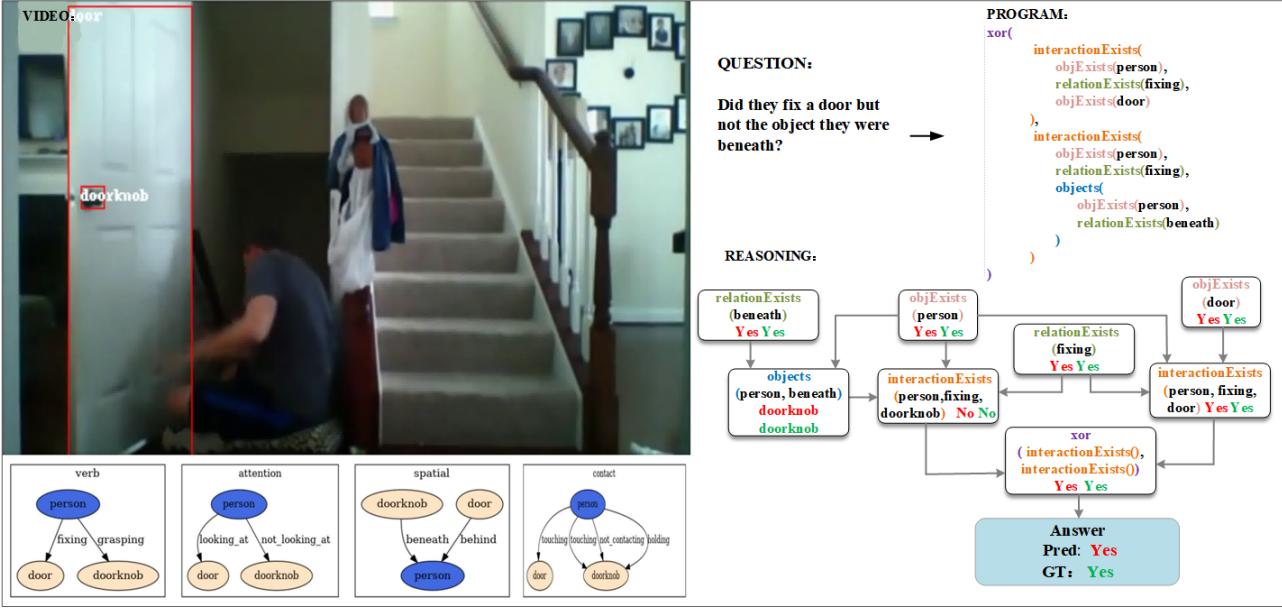


Figure 11. An example of a type *Conjunction* inference process, where red indicates the prediction from NS-VideoQA and green represents the Ground Truth (GT).

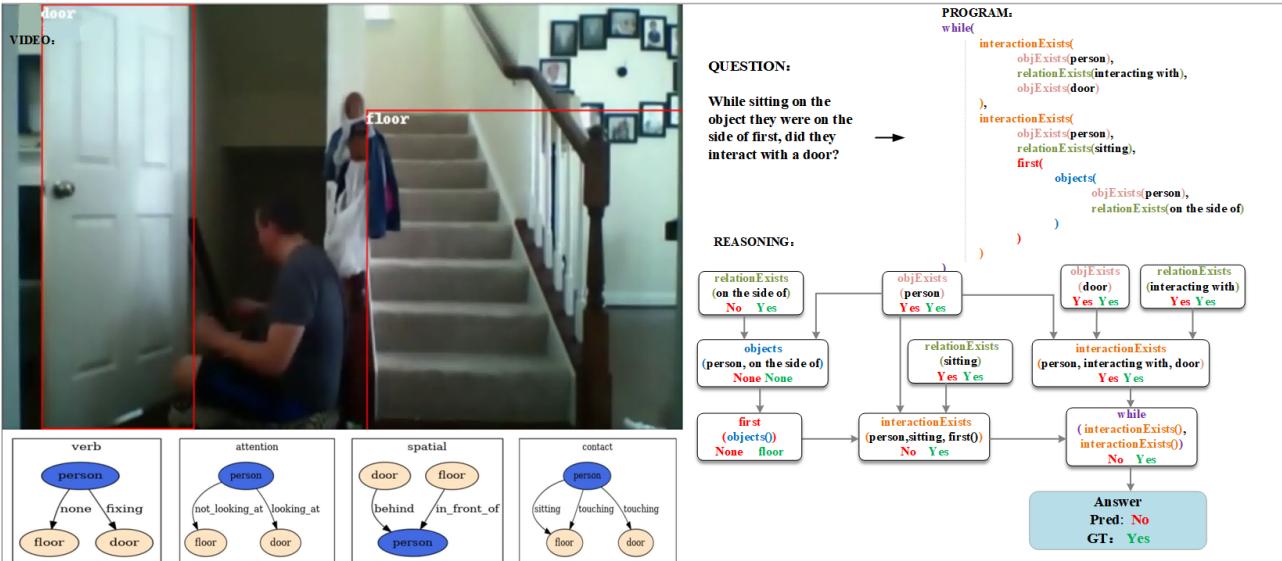


Figure 12. An example of a type *Interaction Temporal Loc.* inference process, where red indicates the prediction from NS-VideoQA and green represents the Ground Truth (GT).

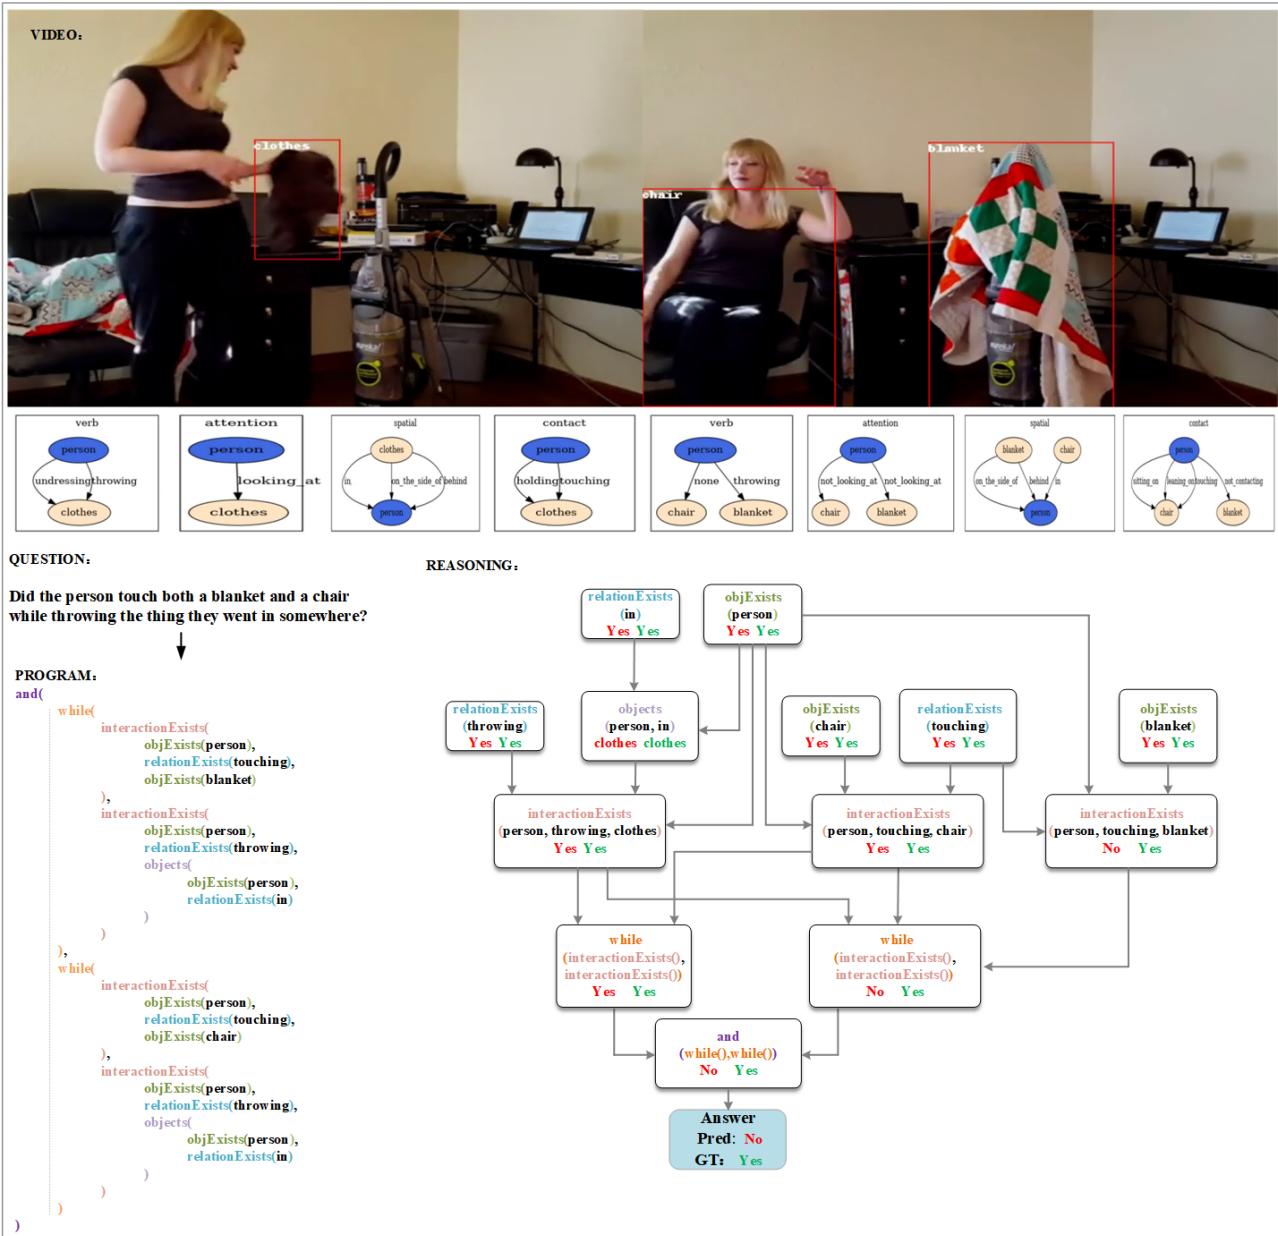


Figure 13. An example of a type *Conjunction* inference process, where red indicates the prediction from NS-VideoQA and green represents the Ground Truth (GT).

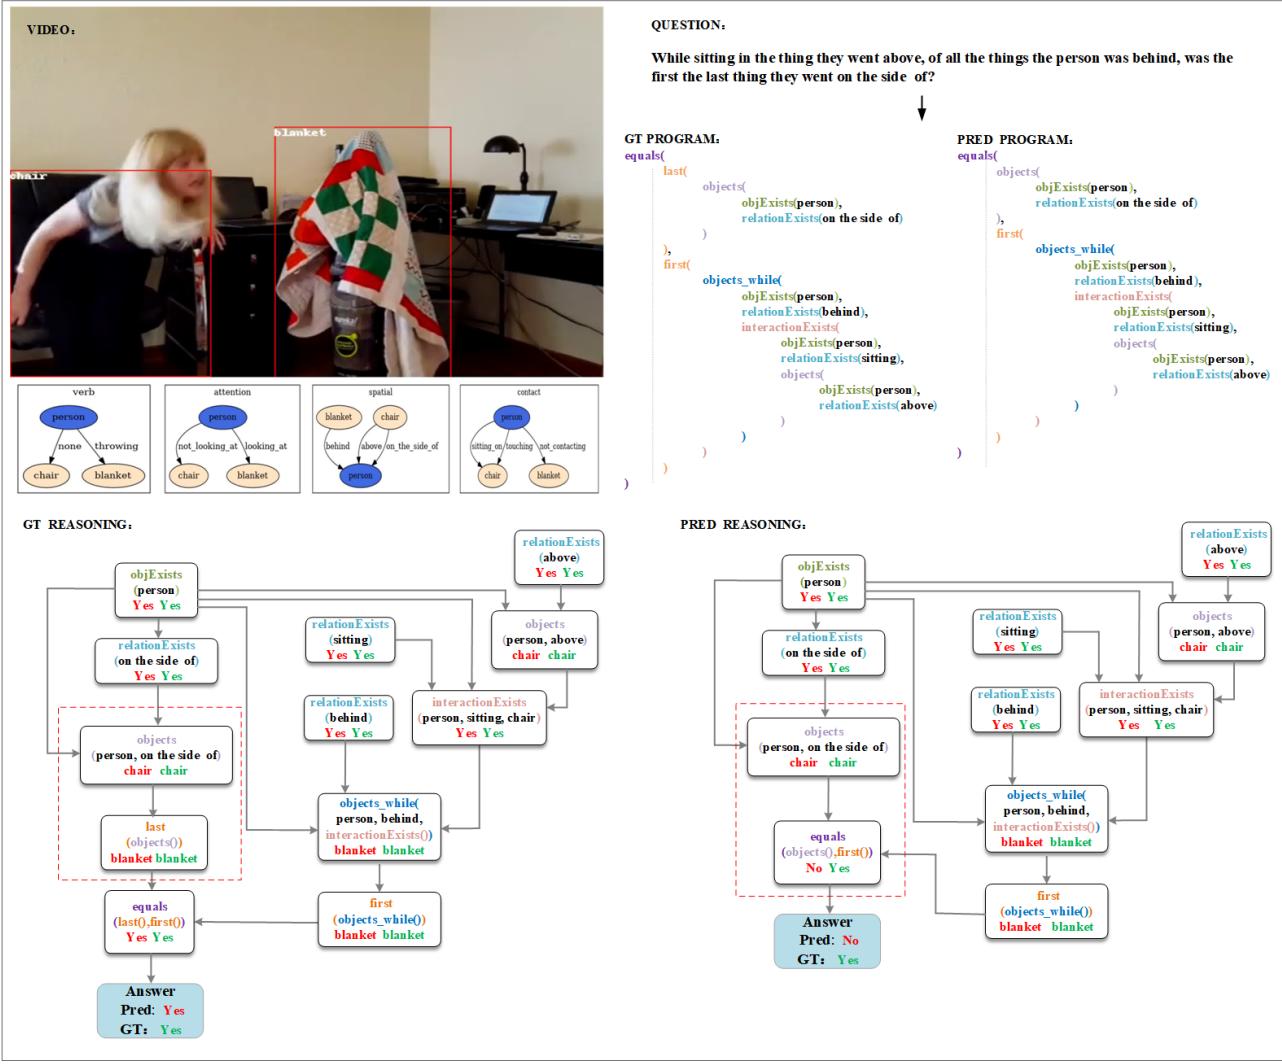


Figure 14. An example of a type *Equals* inference process, where red indicates the prediction from NS-VideoQA and green represents the Ground Truth (GT). The Language Question Parser incorrectly predicts the program `equals (objects (), first())` (marked by the red dotted line in the diagram), leading to a failed final result.

Table 12. Reasoning rules corresponding to *Object Exists*, *Relation Exists*, and *Interaction* question.

Sub-question Type	Reasoning Rule/ Description	Input Type	Output Type
<i>Object Exists</i>	filter_object Selects objects from the input scene list with the input object.	(scene, token)	objects
	query_object To verify if an object exists from objects.		binary
	filter_relation Selects relations from the input scene list with the input relation.	(exe_trace, token)	relations
<i>Relation Exists</i>	query_relation To verify if a relation exists from relations.	(exe_trace, token)	binary
	query_interaction To verify if an interaction exists.	(exe_trace, token)	binary

Table 13. Reasoning rules corresponding to *Temporal loc.*, *Longest/Shortest Action*, *Action*, and *Object* question.

Sub-question Type	Reasoning Rule/ Description	Input Type	Output Type
<i>Interaction Temporal loc./Exists Temporal loc.</i>	interaction_temporal_after Combine two interaction or exists questions using a temporal localizer “after”.	(exe_trace, token)	binary
	interaction_temporal_before Combine two interaction or exists questions using a temporal localizer “before”.	(exe_trace, token)	binary
	interaction_temporal_while Combine two interaction or exists questions using a temporal localizer “while”.	(exe_trace, token)	binary
<i>Action Temporal loc.</i>	interaction_temporal_between Combine two interaction or exists questions using a temporal localizer “between”.		
		(exe_trace, token)	binary
	actions_after Getting the action that occurs after an interaction.		
<i>Object Temporal loc.</i>		(exe_trace, token)	action
	actions_before Getting the action that occurs before an interaction.	(exe_trace, token)	action
<i>Longest/Shortest Action</i>	objects_after Selects objects that involved with the specified interaction appears after the specified action.	(exe_trace, token)	object
	objects_before Selects objects that involved with the specified interaction appears before the specified action.	(exe_trace, token)	object
	objects_while Selects objects that involved with the specified interaction appears while the specified action.	(exe_trace, token)	object
<i>Action</i>	objects_between Selects objects that involved with the specified interaction appears between the two specified actions.		
		(exe_trace, token)	object
<i>Longest/Shortest Action</i>	longest_action Getting the action that takes the longest time to occur.	(exe_trace, token)	action
<i>Action</i>	shortest_action Getting the action that takes the shortest time to occur.	(exe_trace, token)	action
<i>Object</i>	filter_actions Selects actions from the input scene list with the input relation “doing”.	(exe_trace, token)	actions
<i>Object</i>	query_subject_relation Selects objects that involved with the specified interaction.	(exe_trace, token)	object

Table 14. Reasoning rules corresponding to *Choose*, *Equals*, *Conjunction*, and *First/last* question.

Sub-question Type	Reasoning Rule/ Description	Input Type	Output Type
<i>Choose</i>	choose Choose an object of two possible options, such as: “equals”.	(exe_trace, token)	object
	or Choose a word of two possible options, such as: “after” and “before”.		
	choose_action_shorter Choose a shorter action from two possible action.	(exe_trace, token)	action
	choose_action_longer Choose a longer action from two possible action.		
	object.equals Compares two objects to verify if they are the same.	(exe_trace, token)	binary
	action.equals Compares two actions to verify if they are the same.		
<i>Conjunction</i>	conjunction_and Combine two interaction questions using a conjunction “and”.	(exe_trace, token)	binary
	conjunction_xor Combine two interaction questions using a conjunction “xor”.		
<i>First/last</i>	query_first Getting the first instance of the given object or action.	(exe_trace, token)	object/action
	query_last Getting the last instance of the given object or action.		

Table 15. Input/output data types of reasoning rules in the program executor.

Type	Semantics
scene	The structural static and dynamic scene representation: person, objects, relations, actions, and its starting and ending time.
exe_trace	A dictionary that tracks the intermediate result of previous reasoning steps.
token	A string representing the name of the symbolic function to be executed currently.
binary	A string indicating a binary question out of “yes”, “no”.
time	A string indicating a temporal relation out of “after”, “before”.
object	A string indicating an object out of “chair”, “paper”, “food”, “door”, “vacuum”, “person”, “laptop”, “dish”, “phone”, “blanket”, “doorknob”, “clothes”, “window”, “bed”, “floor”, “closet”, “broom”, “mirror”, “table”, “refrigerator”, “pillow”, “picture”, “bag”, “box”, “light”, “shoe”, “medicine”, “doorway”, “television”.
objects	A list of the specified object in the scene.
relation	A string indicating a relation out of “looking at”, “not looking at”, “unseen”, “above”, “beneath”, “in front of”, “behind”, “on the side of”, “in”, “carrying”, “covered by”, “drinking from”, “eating”, “having it on the back”, “holding”, “leaning on”, “lying on”, “sitting on”, “standing on”, “touching”, “twisting”, “wearing”, “wiping”, “writing on”, “not contacting”, “drinking”, “putting”, “taking”, “closing”, “throwing”, “putting down”, “grasping”, “walking”, “sitting”, “watching”, “opening”, “snuggling”, “standing”, “working on”, “tidying”, “working”, “awakening”, “fixing”, “smiling”, “playing”, “lying”, “playing on”, “sneezing”, “dressing”, “undressing”, “washing”, “pouring”, “turning”, “making”, “going”, “talking”, “consuming”, “laughing”, “running”, “reaching”, “photographing”, “cooking”.
relations	A list of the specified relation in the scene.
action	A string indicating an action out of “undressing themselves”, “fixing a vacuum”, “washing a mirror”, “holding a bag”, “snuggling with a pillow”, “watching a picture”, “watching a laptop or something on a laptop”, “fixing a door”, “holding a vacuum”, “putting on a shoe”, “holding some food”, “washing something with a blanket”, “watching a book”, “turning off a light”, “holding a blanket”, “watching television”, “holding a mirror”, “taking off some shoes”, “sitting at a table”, “washing a window”, “fixing their hair”, “fixing a doorknob”, “tidying up a blanket”, “holding a book”, “washing a cup”, “lying on the floor”, “tidying up with a broom”, “holding a paper”, “smiling at something”, “working on a book”, “holding a broom”, “holding a cup of something”, “watching something in a mirror”, “holding some medicine”, “laughing at something”, “fixing a light”, “snuggling with a blanket”, “holding some clothes”, “holding a phone”, “washing some clothes”, “holding a picture”, “pouring something into a cup”, “dressing themselves”, “tidying up a closet”, “sitting in a bed”, “holding a shoe”, “holding a pillow”, “washing their hands”, “None”, “turning on a light”, “lying on a bed”, “tidying some clothes”, “washing a table”, “tidying something on the floor”, “sitting on the floor”, “tidying up a table”, “standing up”, “walking through a doorway”, “eating some food”, “holding a dish”, “standing on a chair”, “watching outside of a window”, “grasping onto a doorknob”, “holding a box”, “running somewhere”, “sitting in a chair”, “holding a laptop”, “making some food”, “sitting on a table”, “awakening in bed”, “sneezing somewhere”.
actions	A list of the specified relation “doing” in the scene.